



UNIVERSITY OF COPENHAGEN

Network Apps: Overview of Socket API

Network Apps: HTTP & Content Delivery

Vivek Shah

Based on slides compiled by Marcos Vaz Salles

Recap: Numerical example

- How long does it take to send a file of 640,000 bits from host A to host B over a circuit-switched network?
 - all link speeds: 1536 Mbps
 - each link uses TDM with 24 slots/sec
 - 500 msec to establish end-to-end circuit
 - Note: 1 Mbps = 10^6 bps
- Possible answers
 - (a) 500 msec
 - (b) 500.4 msec
 - (c) 510 msec
 - (d) 1 sec



Recap: Key Concepts in Networking

- **Protocols**
 - Speaking the same language
 - Syntax and semantics
- **Layering**
 - Standing on the shoulders of giants
 - A key to managing complexity
- **Resource allocation**
 - Dividing scarce resources among competing parties
 - Memory, link bandwidth, wireless spectrum, paths
- **Naming**
 - What to call computers, services, protocols, ...



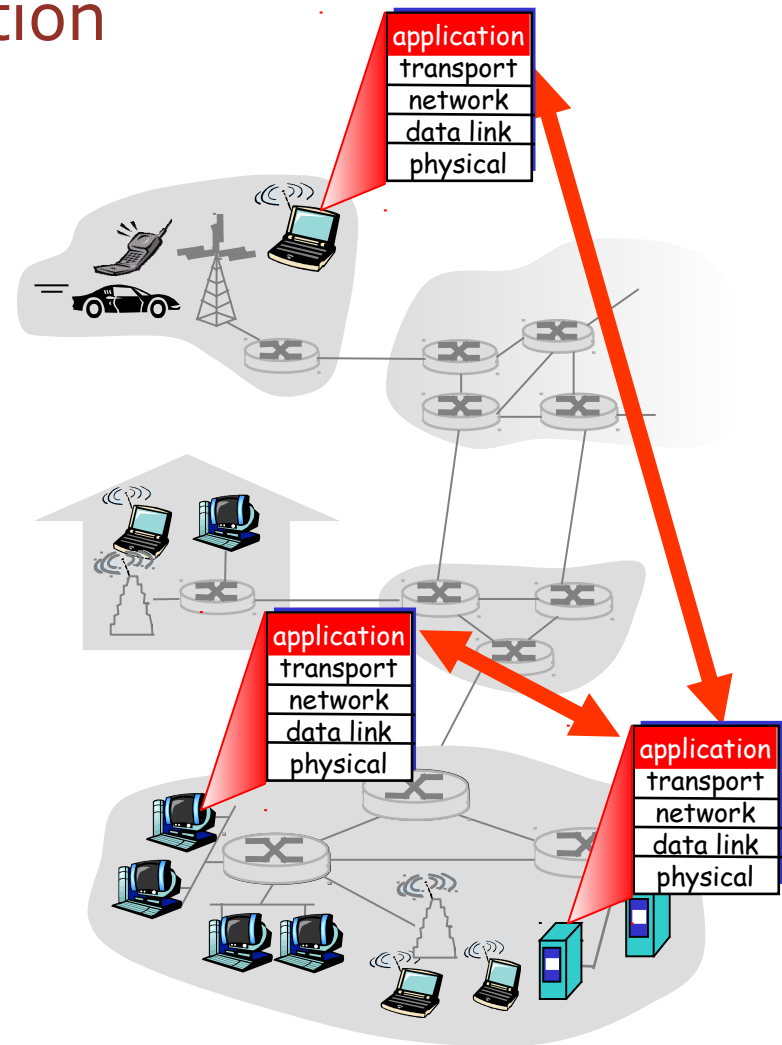
Creating a network application

write programs that

- run on (different) *end systems*
- communicate over network
- e.g., web server software communicates with browser software

No need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation



Source: Kurose & Ross


Some network apps

- e-mail
- remote login
- web
- instant messaging
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube)
- voice over IP
- real-time video conferencing
- social networking
- cloud computing
- ...
- ...
- ...



Network applications need streams of data

Circuit switching Packet switching



http://www.tcpipguide.com/free/t_CircuitSwitchingandPacketSwitchingNetworks-2.htm

Today's networks provide packet delivery, not streams!



Source: Freedman (partial)

What if the Data Doesn't Fit?

```
GET /courses/archive/spr09/cos461/ HTTP/1.1
```

```
Host: www.cs.princeton.edu
```

```
User-Agent: Mozilla/4.03
```

```
CRLF
```


Request



Problem: Packet size

- Typical Web page is 10 kbytes
- On Ethernet, max IP packet is 1500 bytes

GET index.html



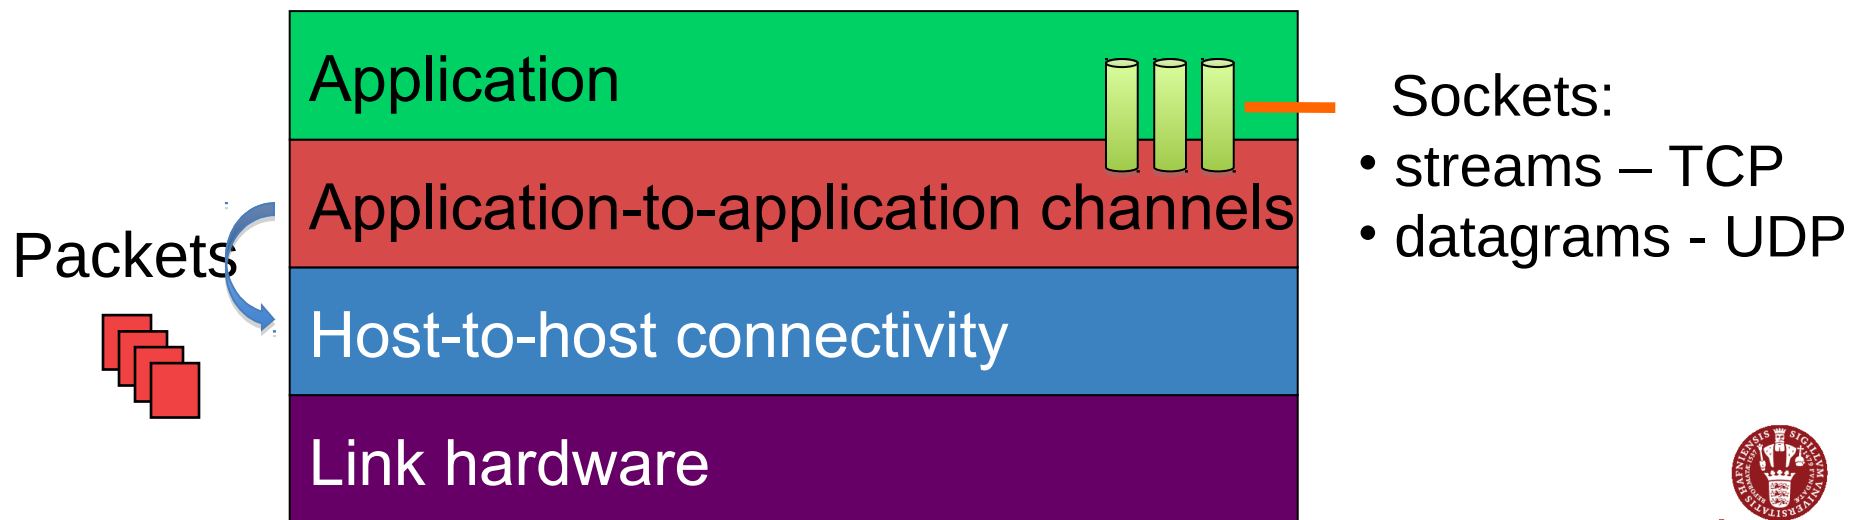
Solution: Split the data across multiple packets

Source: Freedman



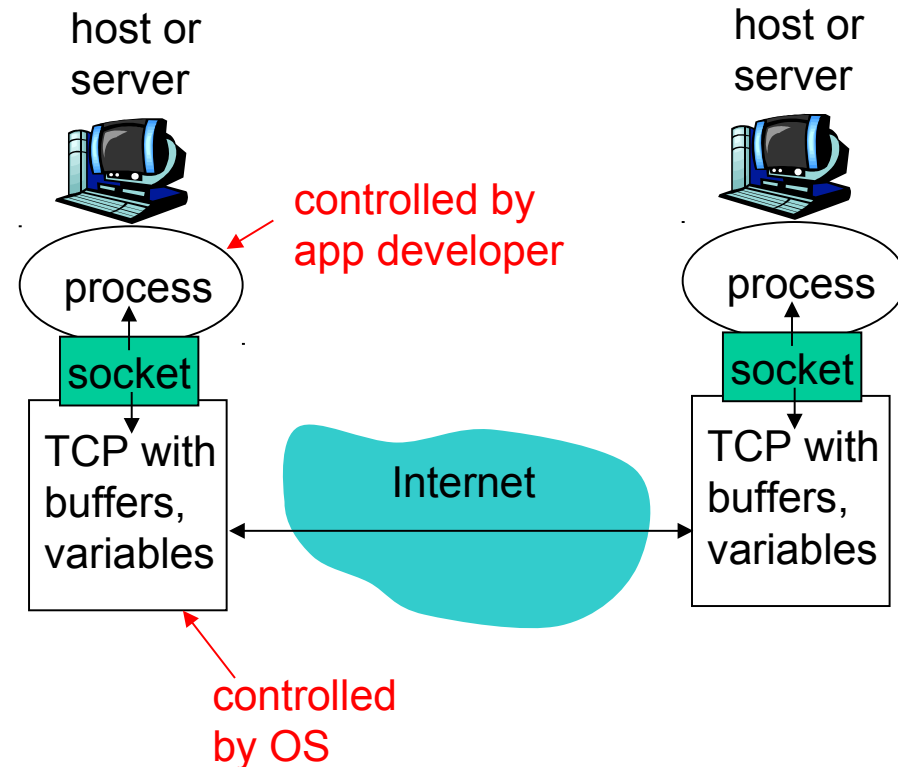
Layering = Functional Abstraction

- Sub-divide the problem
 - Each layer relies on services from layer below
 - Each layer exports services to layer above
- Interface between layers defines interaction
 - Hides implementation details
 - Layers can change without disturbing other layers



Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door which brings message to socket at receiving process



API: (1) choice of transport protocol; (2) ability to fix a few parameters (more on this in next lecture!)

Internet transport protocols services

TCP service:

- *connection-oriented*: setup required between client and server processes
- *reliable transport*: between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantees, security

UDP service:

- unreliable data transfer between sending and receiving process
- does not provide: connection setup, reliability, flow control, congestion control, timing, throughput guarantee, or security

Q: why bother? Why is there a UDP?



UNIX Socket API

- Socket interface
 - Originally provided in Berkeley UNIX
 - Later adopted by all popular operating systems
 - Simplifies porting applications to different OSes
- In UNIX, everything is like a file
 - All input is like reading a file
 - All output is like writing a file
 - File is represented by an integer file descriptor
- API implemented as system calls
 - E.g., connect, read, write, close, ...

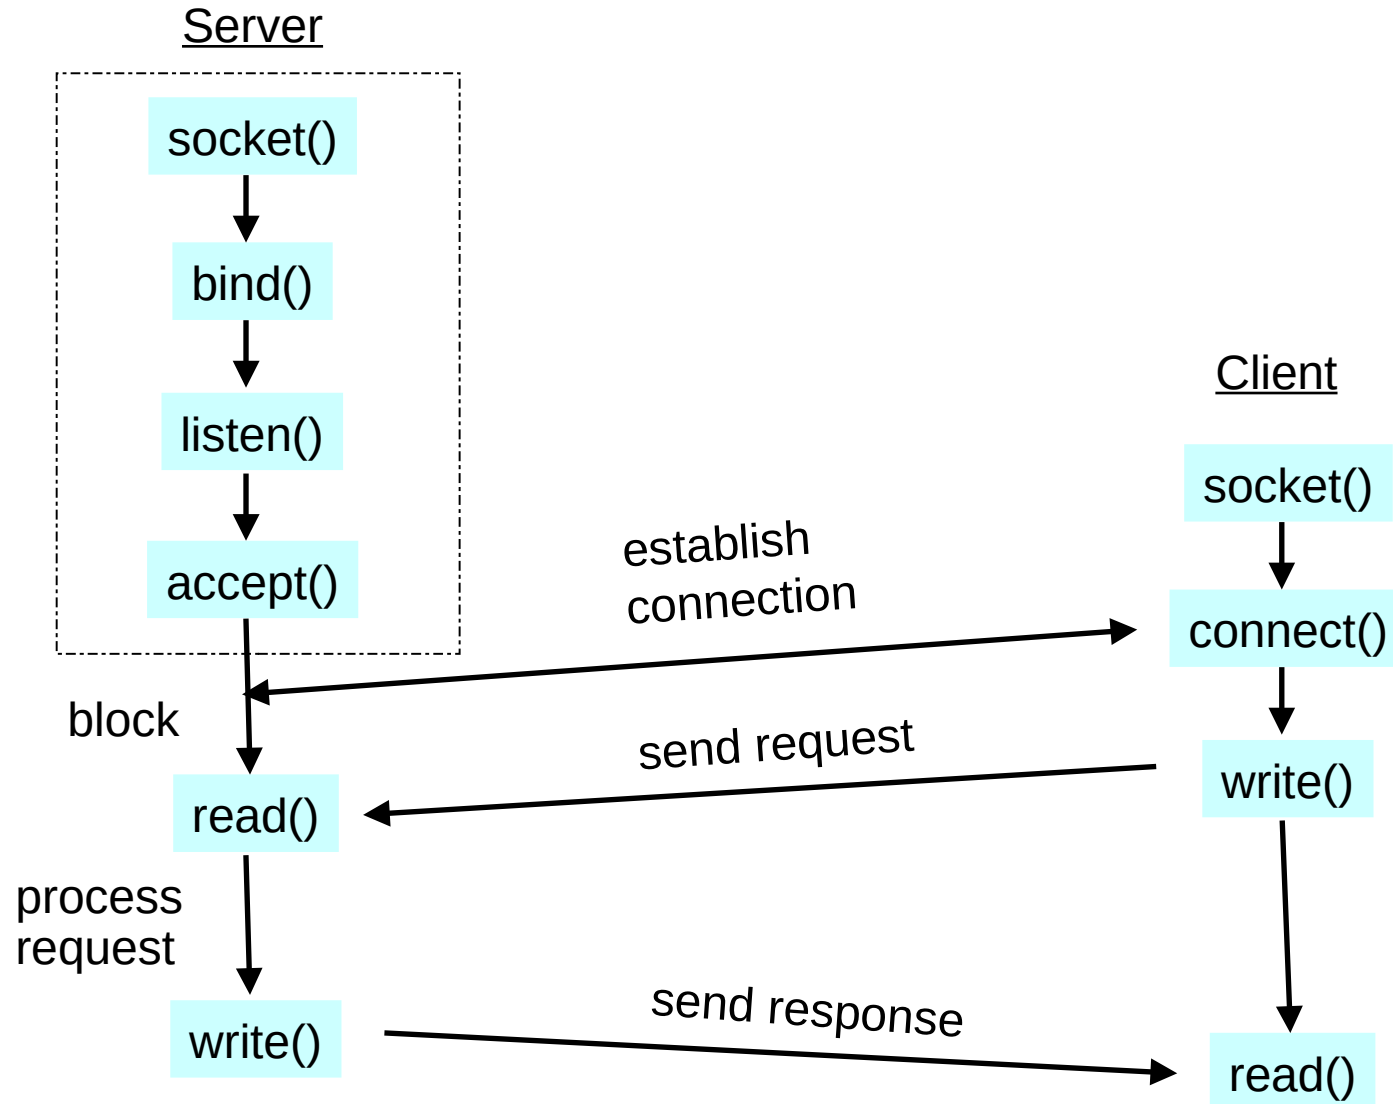


Identifying the Receiving Process

- Sending process must identify the receiver
 - The receiving end host machine
 - The specific socket in a process on that machine
- Receiving host
 - Destination address that uniquely identifies the host
 - Typically, high-level name translated to IP address (DNS)
 - For example, www.diku.dk → **130.225.96.108**
 - An IP address is a 32-bit quantity
- Receiving socket
 - Host may be running many different processes
 - Destination port that uniquely identifies the socket
 - A port number is a 16-bit quantity




Client-Server TCP Sockets



Source:
Freedman
(partial)


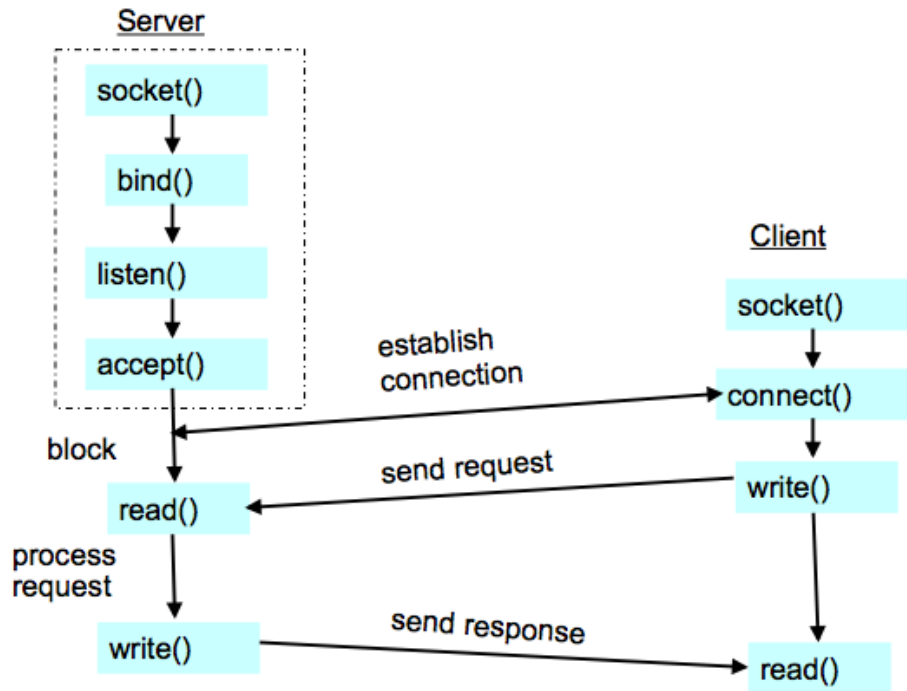


Application Architectures



- P2P highly scalable, but difficult to manage
- Hybrids also possible, e.g., Skype

Discussion: How do you set sockets up in a P2P application?




- How many sockets in each peer?
- What if many peers connect to one peer?

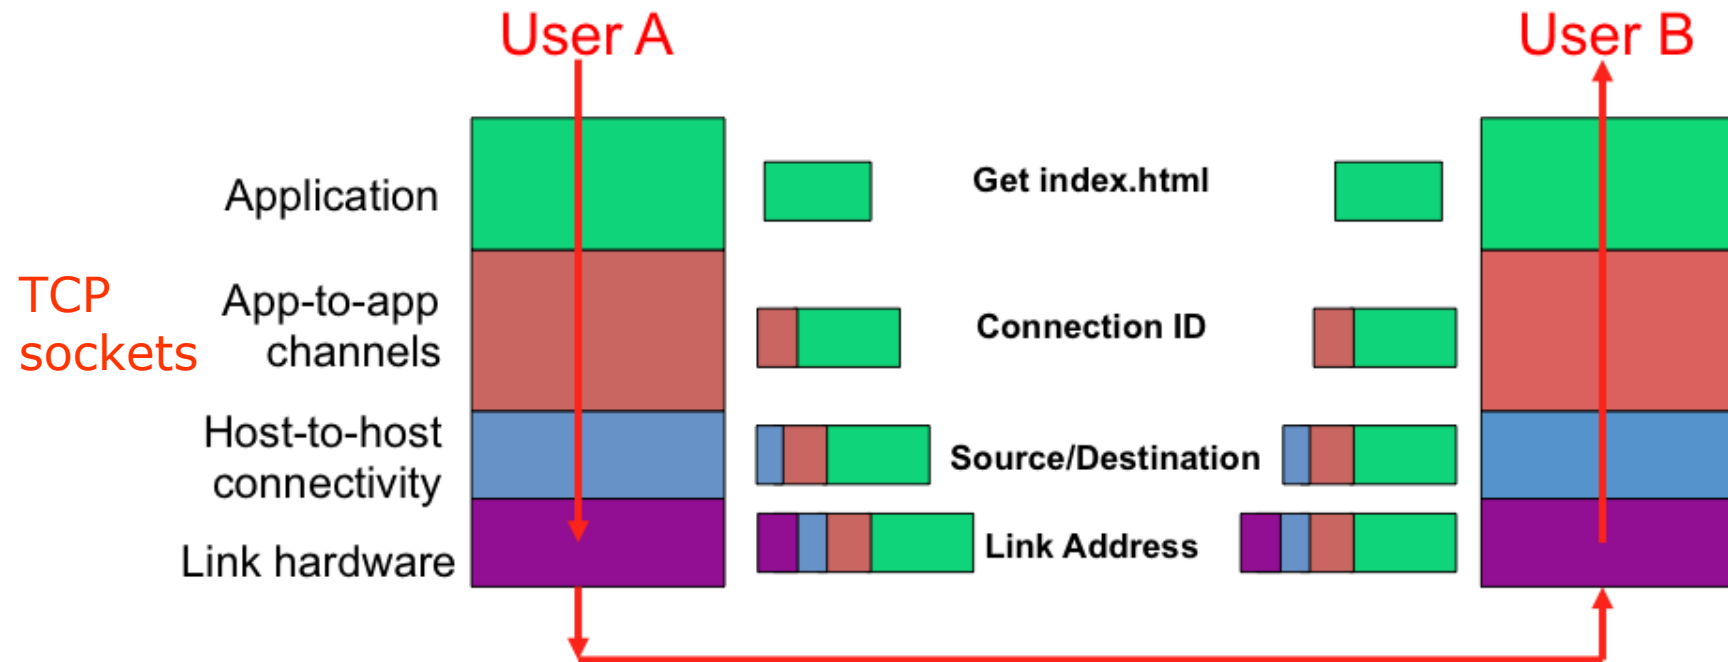
HTTP Basics

- HTTP layered over bidirectional byte stream
- Interaction
 - Client sends request to server, followed by response from server to client
 - Requests/responses are encoded in text
- Targets access to web objects
 - GET, POST, HEAD → HTTP/1.0
 - GET, POST, HEAD, PUT, DELETE → HTTP/1.1
- Stateless
 - Server maintains no info about past client requests
 - What about personalization? Data stored in back-end database; client sends “web cookie” used to lookup data





Layer Encapsulation in HTTP



HTTP Request Example

GET / HTTP/1.1

Host: sns.cs.princeton.edu

Accept: */*

Accept-Language: en-us

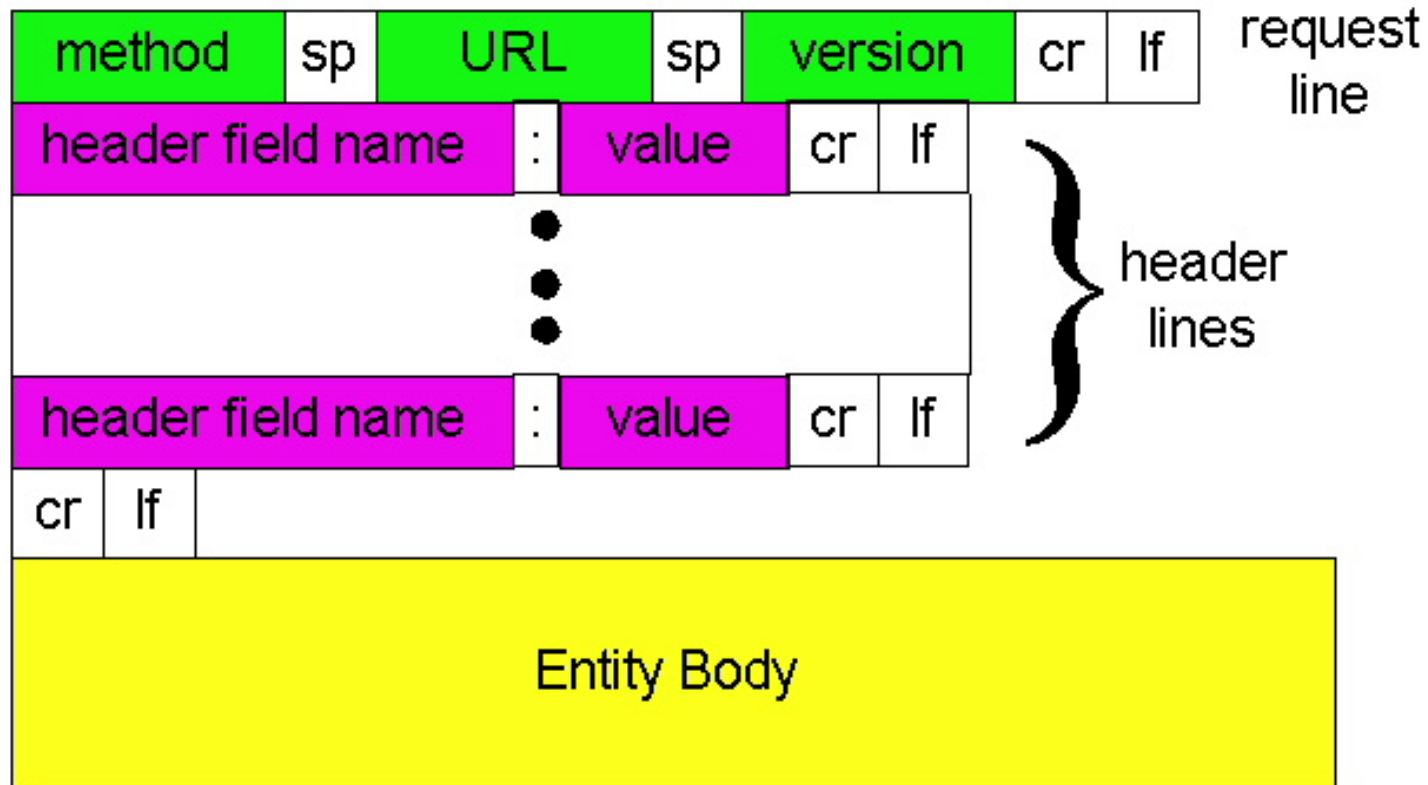
Accept-Encoding: gzip, deflate

User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X 10.5;
en-US; rv:1.9.2.13) Gecko/20101203 Firefox/3.6.13

Connection: Keep-Alive



HTTP Request



HTTP Response Example

HTTP/1.1 200 OK

Date: Wed, 02 Feb 2011 04:01:21 GMT

Server: Apache/2.2.3 (CentOS)

X-Pingback: <http://sns.cs.princeton.edu/xmlrpc.php>

Last-Modified: Wed, 01 Feb 2011 12:41:51 GMT

ETag: "7a11f-10ed-3a75ae4a"

Accept-Ranges: bytes

Content-Length: 4333

Keep-Alive: timeout=15, max=100

Connection: Keep-Alive

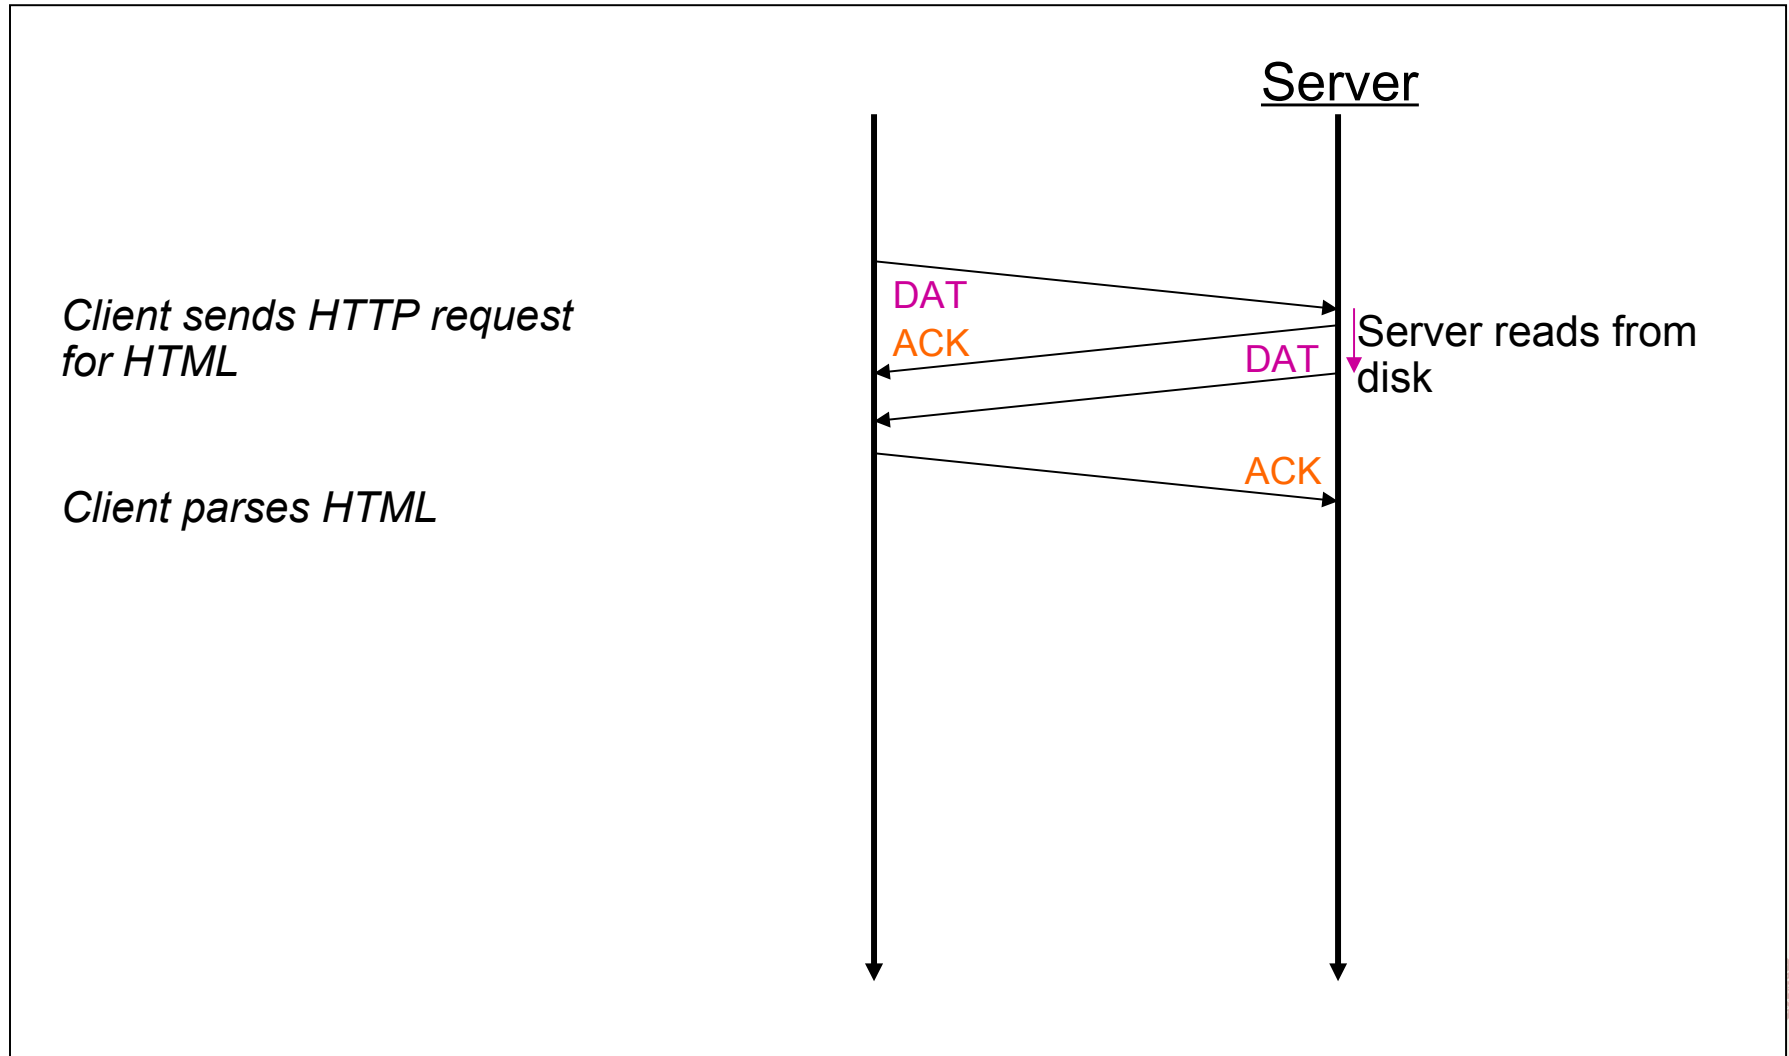
```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN"  
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
```

```
<html xmlns="http://www.w3.org/1999/xhtml" dir="ltr" lang="en-US">
```



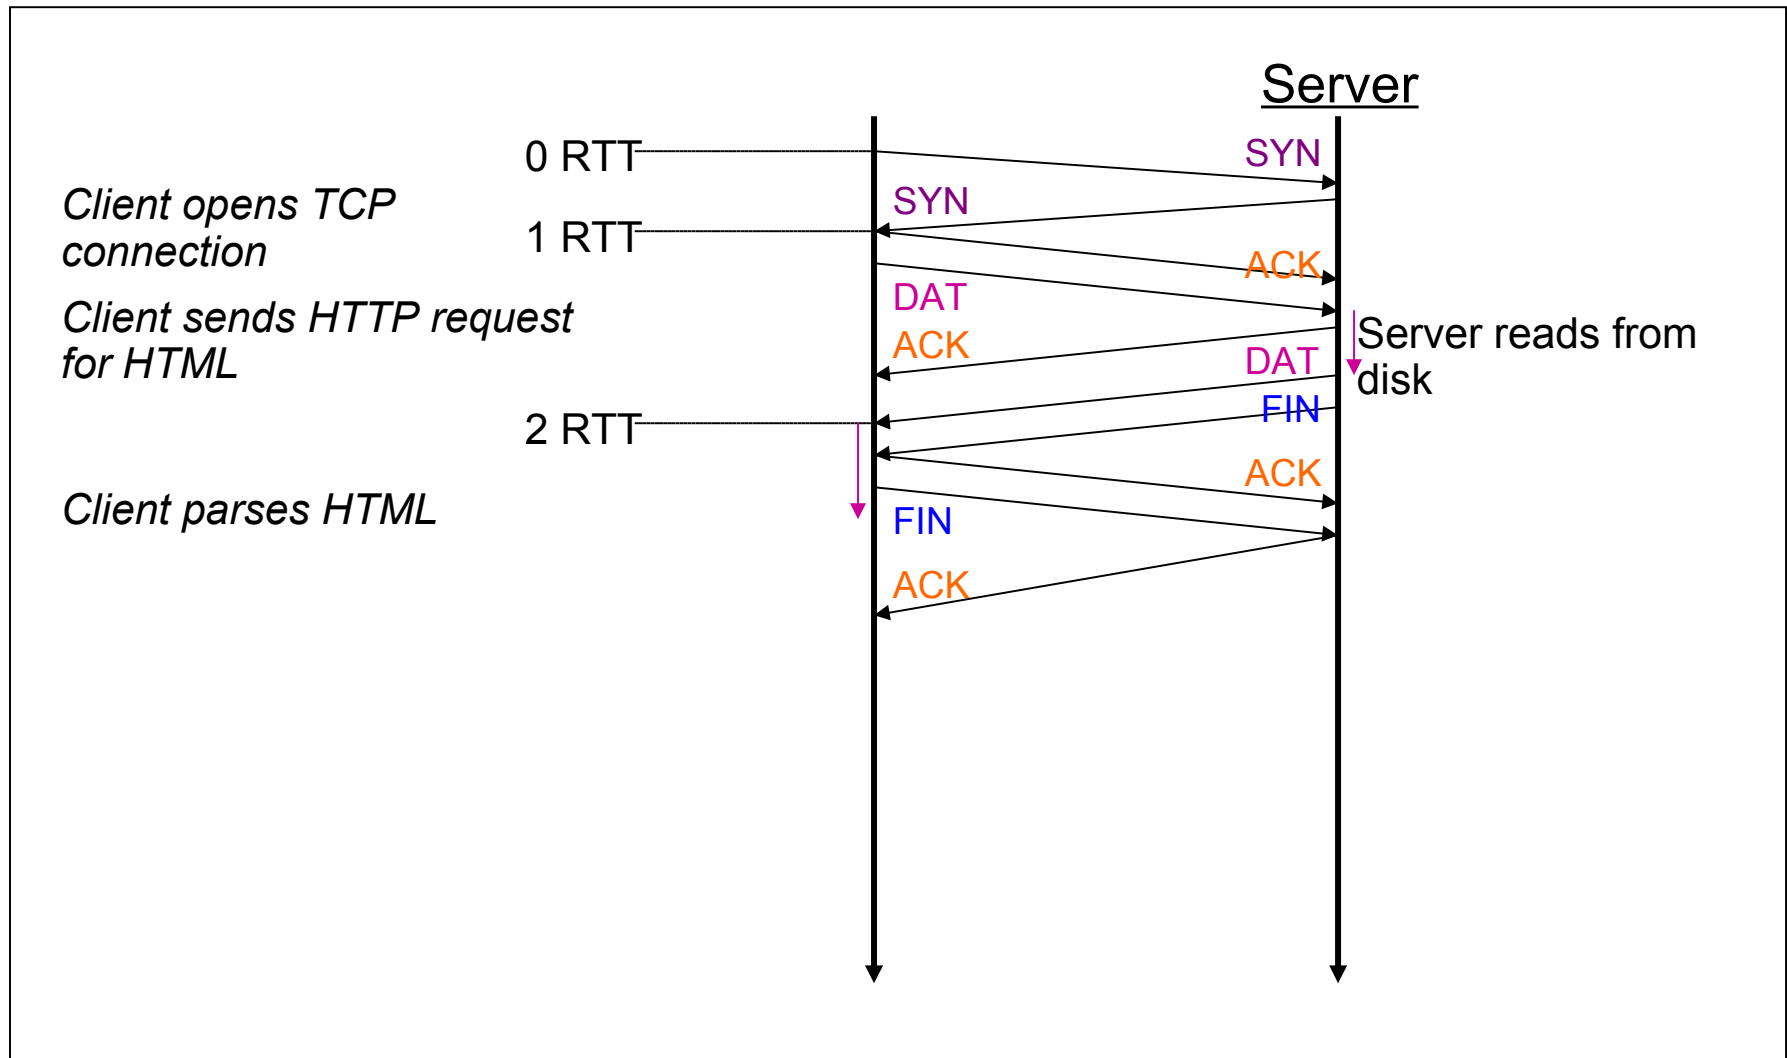
Single Transfer Example

Source: Freedman



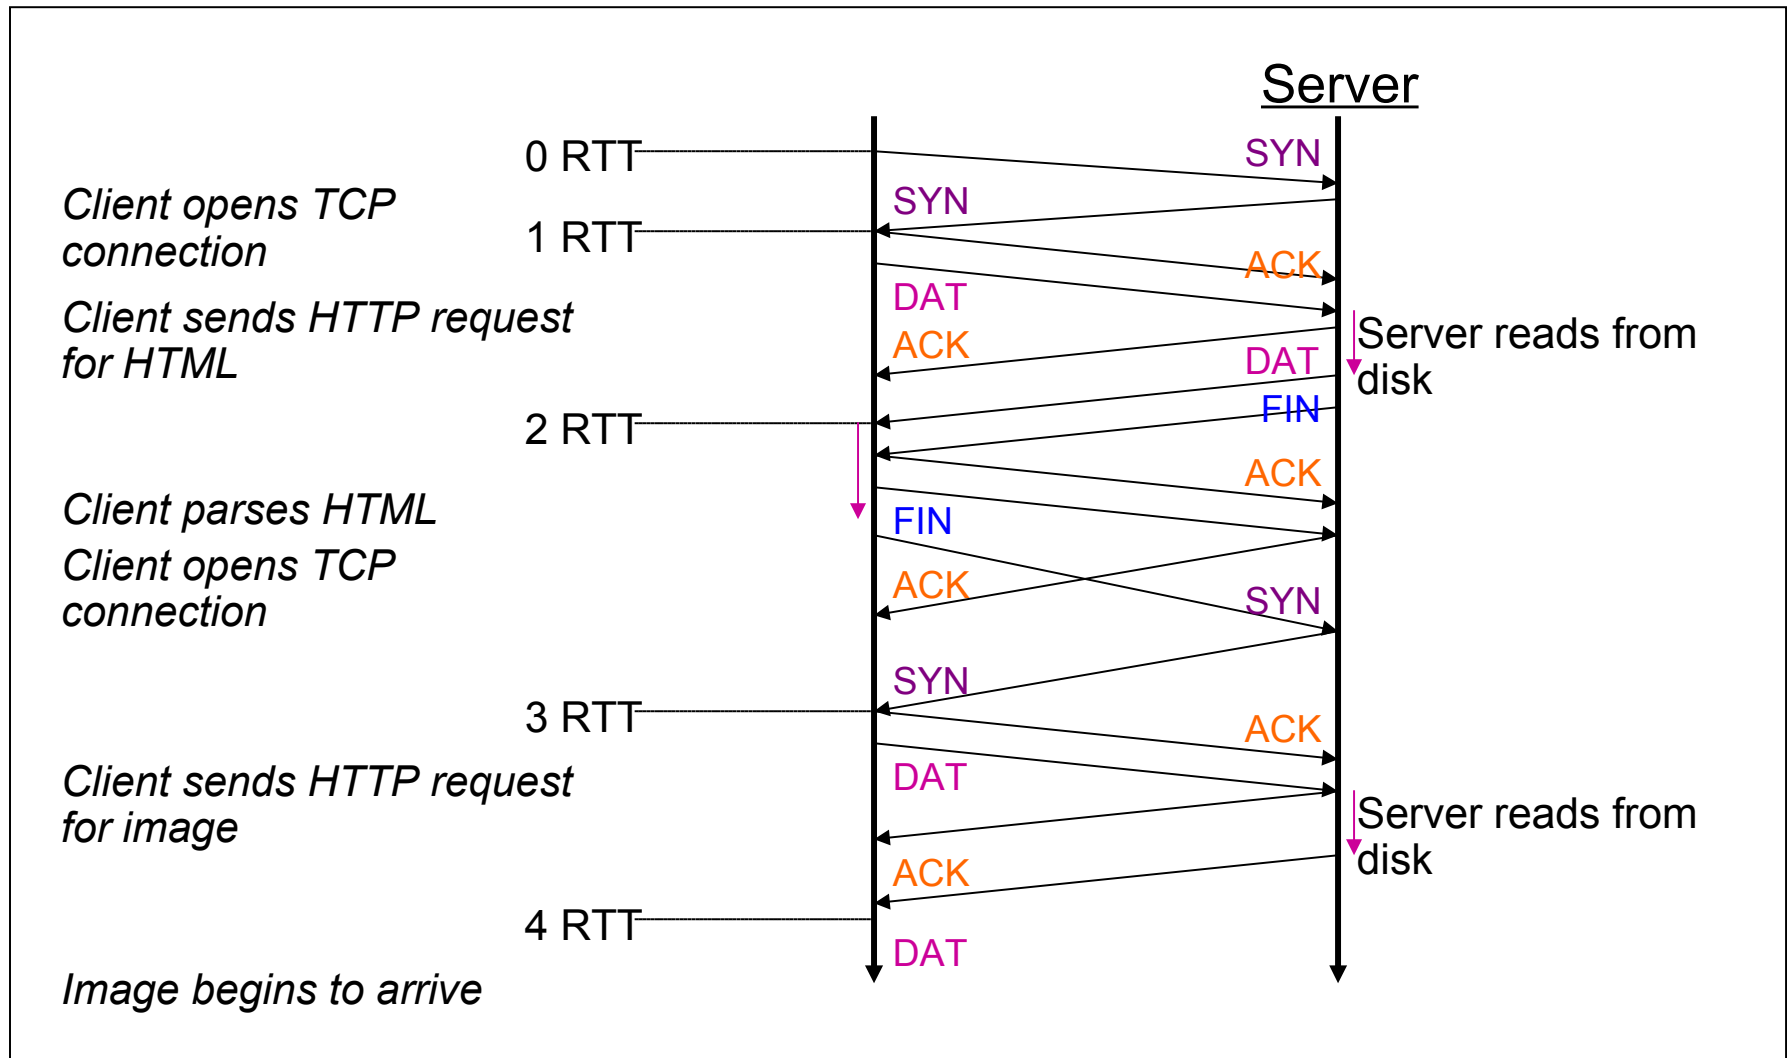
Single Transfer Example

Source: Freedman



Single Transfer Example

Source: Freedman



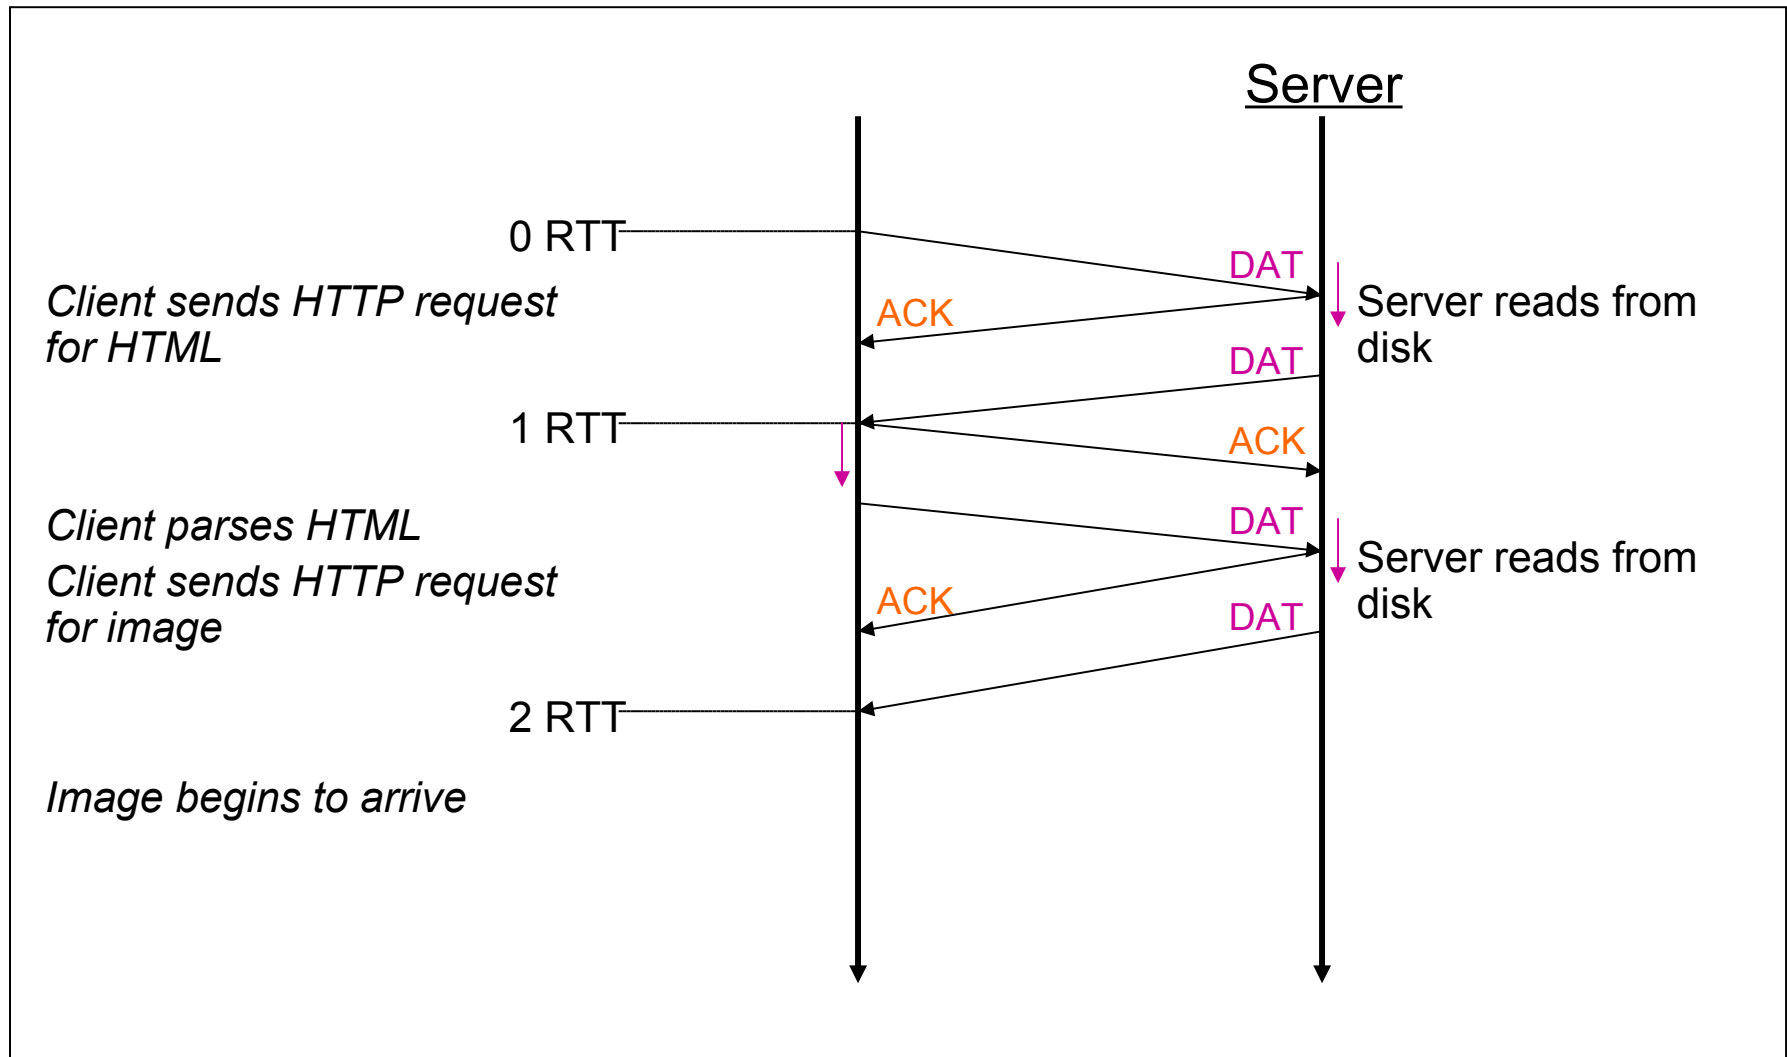
Problems with simple model

- Multiple connection setups
 - Three-way handshake each time (TCP “synchronizing” stream)
- Lots of extra connections
 - Increases server state/processing
 - Server forced to keep connection state
- Later we will see also that
 - Short transfers are hard on stream protocol (TCP)
 - How much data should it send at once?
 - Congestion avoidance: Takes a while to “ramp up” to high sending rate (TCP “slow start”)
 - Loss recovery is poor when not “ramped up”



Persistent Connection Example

Source: Freedman



Persistent HTTP

Non-persistent HTTP issues:

- Requires 2 RTTs per object
- OS must allocate resources for each TCP connection
- But browsers often open parallel TCP connections to fetch referenced objects

Persistent HTTP:

- Server leaves connection open after sending response
- Subsequent HTTP messages between same client/server are sent over connection



Persistent HTTP

Persistent without pipelining:

- Client issues new request only when previous response has been received
- One RTT for each object

Persistent with pipelining:

- Default in HTTP/1.1 spec
- Client sends requests as soon as it encounters referenced object
- As little as one RTT for all the referenced objects
- Server must handle responses in same order as requests

• Persistent without pipelining most common:
When does pipelining work best?

• Multiple parallel requests or pipelined requests?

Source: Freedman (partial)



HTTP Caching

- Clients often cache documents
 - When should origin be checked for changes?
 - Every time? Every session? Date?
- HTTP includes caching information in headers
 - HTTP 0.9/1.0 used: "Expires: <date>"; "Pragma: no-cache"
 - HTTP/1.1 has "Cache-Control"
 - "No-Cache", "Private", "Max-age: <seconds>"
 - "E-tag: <opaque value>"
- If not expired, use cached copy
- If expired, use condition GET request to origin
 - "If-Modified-Since: <date>", "If-None-Match: <etag>"
 - 304 ("Not Modified") or 200 ("OK") response



HTTP Conditional Request

GET / HTTP/1.1

Host: sns.cs.princeton.edu

User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac
OS X 10.5; en-US; rv:1.9.2.13)

Connection: Keep-Alive

If-Modified-Since: Tue, 1 Feb 2011 17:54:18 GMT

If-None-Match: "7a11f-10ed-3a75ae4a"

HTTP/1.1 304 Not Modified

Date: Wed, 02 Feb 2011 04:01:21
GMT

Server: Apache/2.2.3 (CentOS)

ETag: "7a11f-10ed-3a75ae4a"

Accept-Ranges: bytes


Keep-Alive: timeout=15, max=100

Connection: Keep-Alive




Web Proxy Caches

- User configures browser: Web accesses via cache
- Browser sends all HTTP requests to cache
 - Object in cache: cache returns object
 - Else: cache requests object from origin, then returns to client




When a single cache isn't enough

- What if the working set is $>$ proxy disk?
 - Cooperation!
- A static hierarchy
 - Check local
 - If miss, check siblings
 - If miss, fetch through parent
- Internet Cache Protocol (ICP)
 - ICPv2 in RFC 2186 (& 2187)
 - UDP-based, short timeout



Web traffic has cacheable workload



“Zipf” or “power-law” distribution

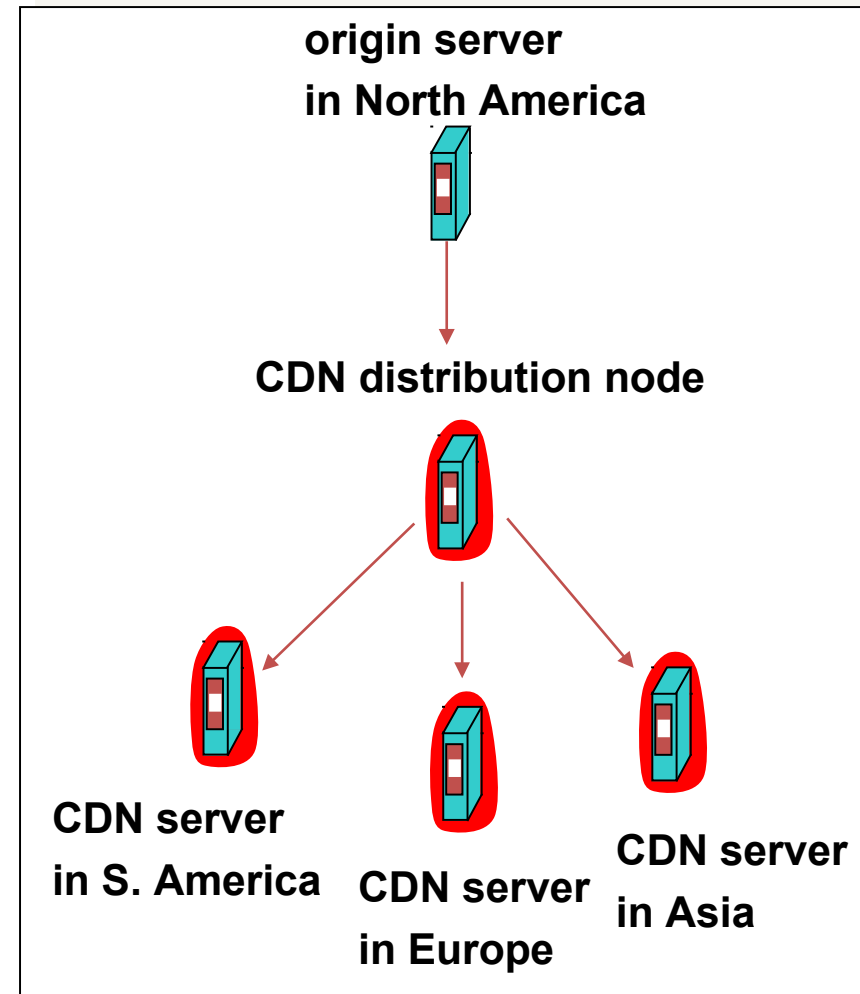
Source: Freedman

Content Distribution Networks (CDNs)

- Content providers are CDN customers

Content replication

- CDN company installs thousands of servers throughout Internet
 - In large datacenters
 - Or, close to users
- CDN replicates customers' content
- When provider updates content, CDN updates servers



Content Distribution Networks & Server Selection

- Replicate content on many servers
- Challenges
 - How to replicate content
 - Where to replicate content
 - How to find replicated content
 - How to choose among known replicas
 - How to direct clients towards replica



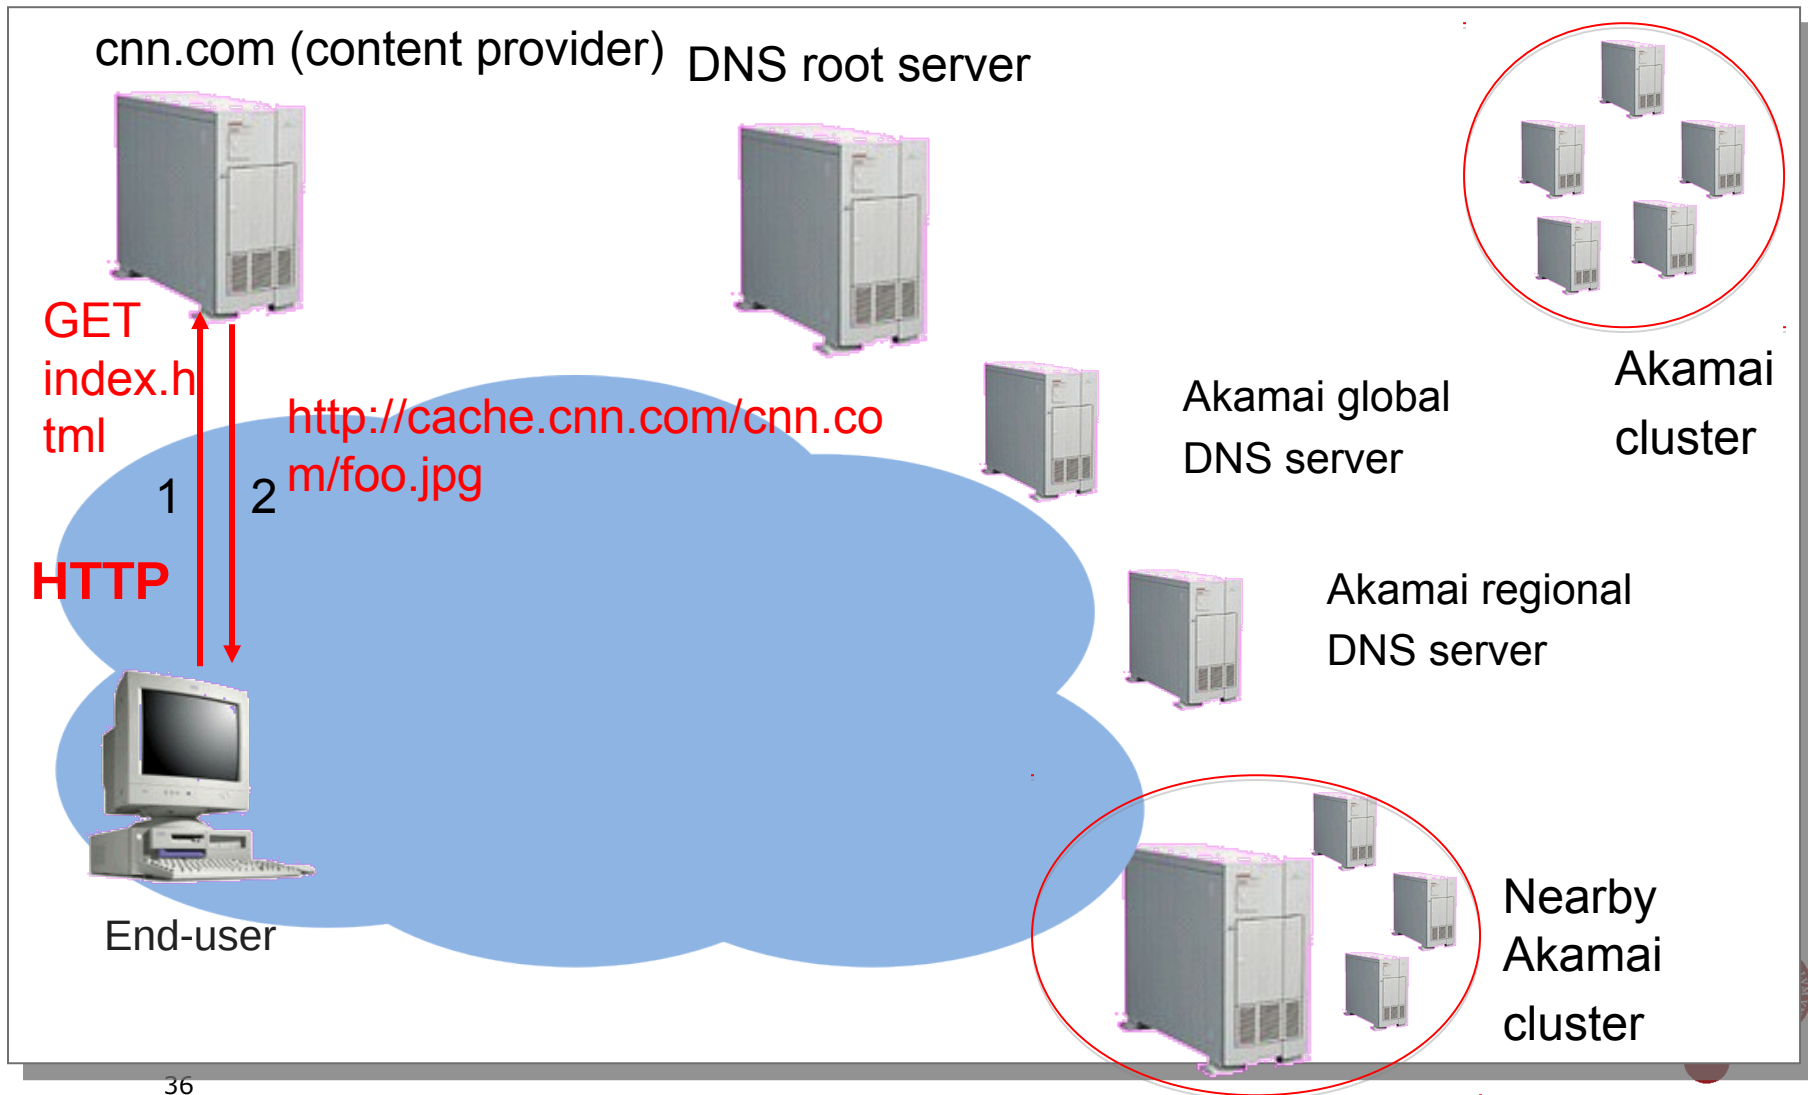
Server Selection

- Which server?
 - Lowest load: to balance load on servers
 - Best performance: to improve client performance
 - Based on Geography? RTT? Throughput? Load?
 - Any alive node: to provide fault tolerance
- How to direct clients to a particular server?
 - As part of routing: anycast, cluster load balancing
 - As part of application: HTTP redirect
 - As part of naming: DNS
- We will explain some of these techniques better later in the course!



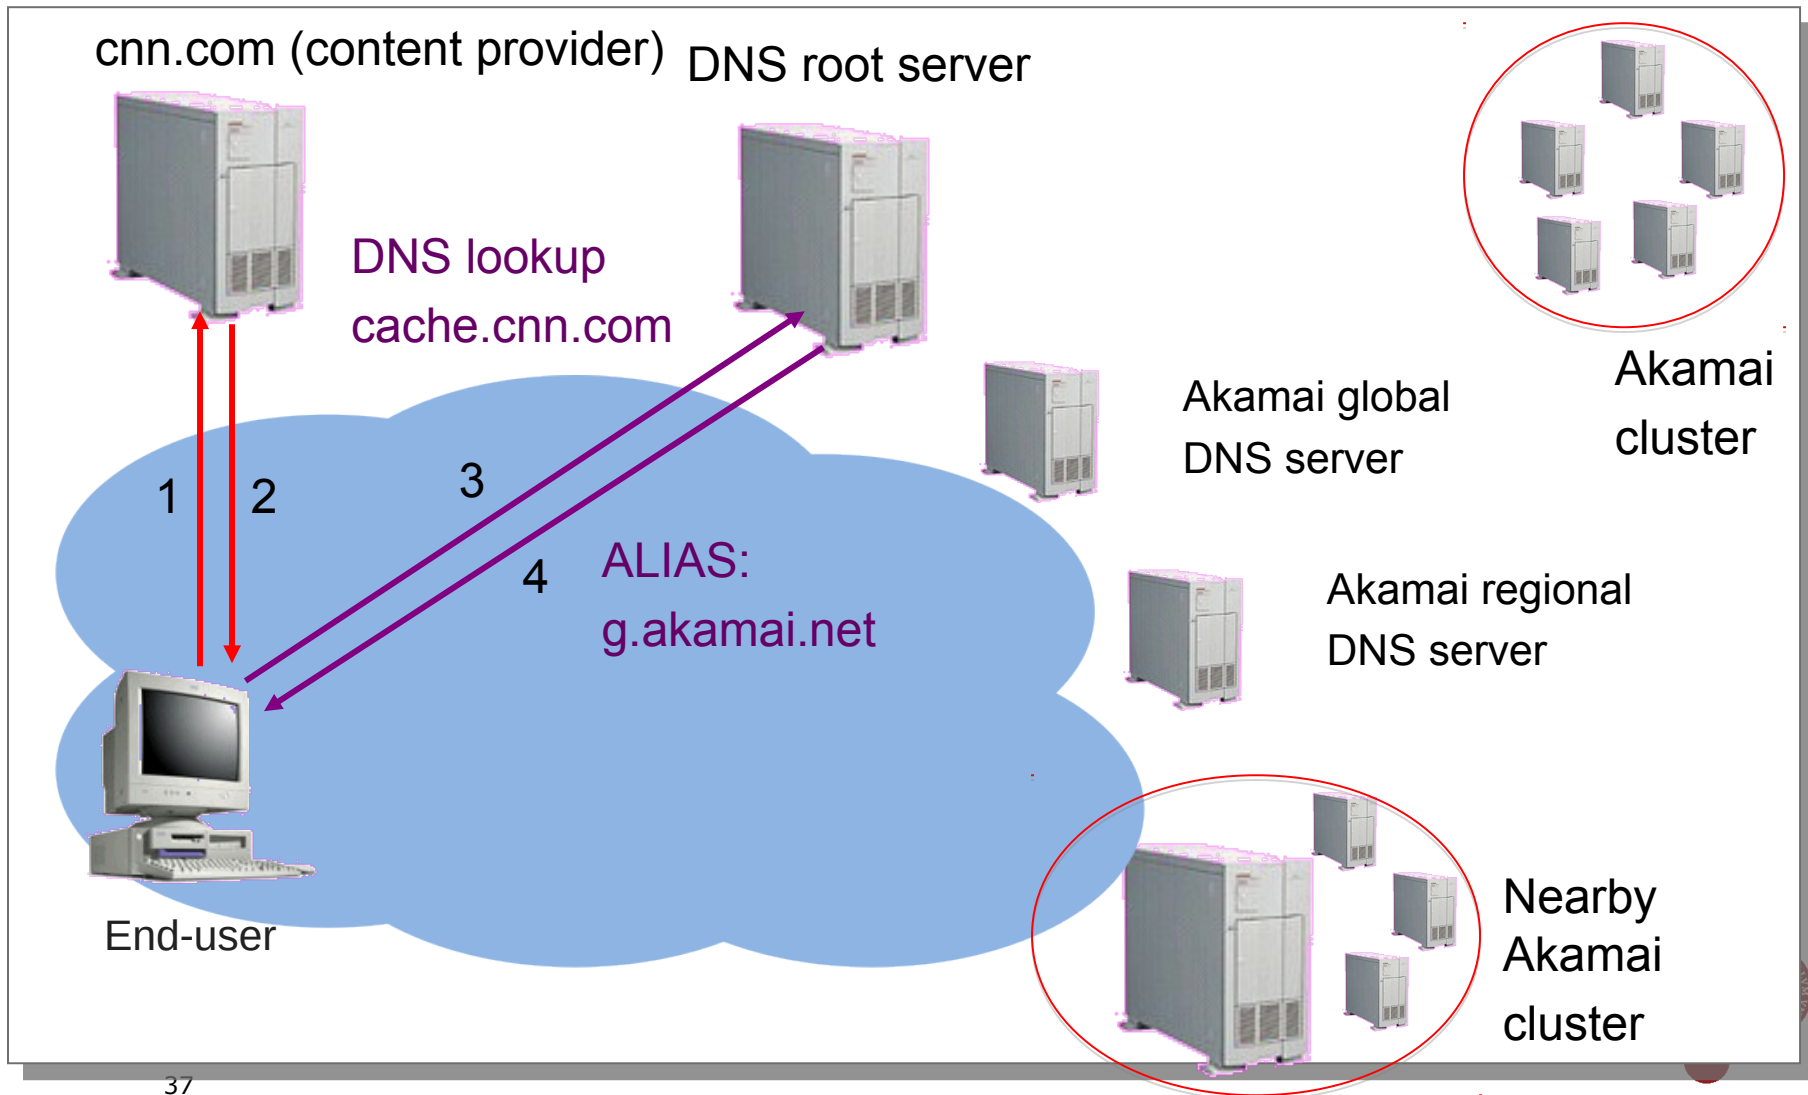
How Akamai Works

Source: Freedman



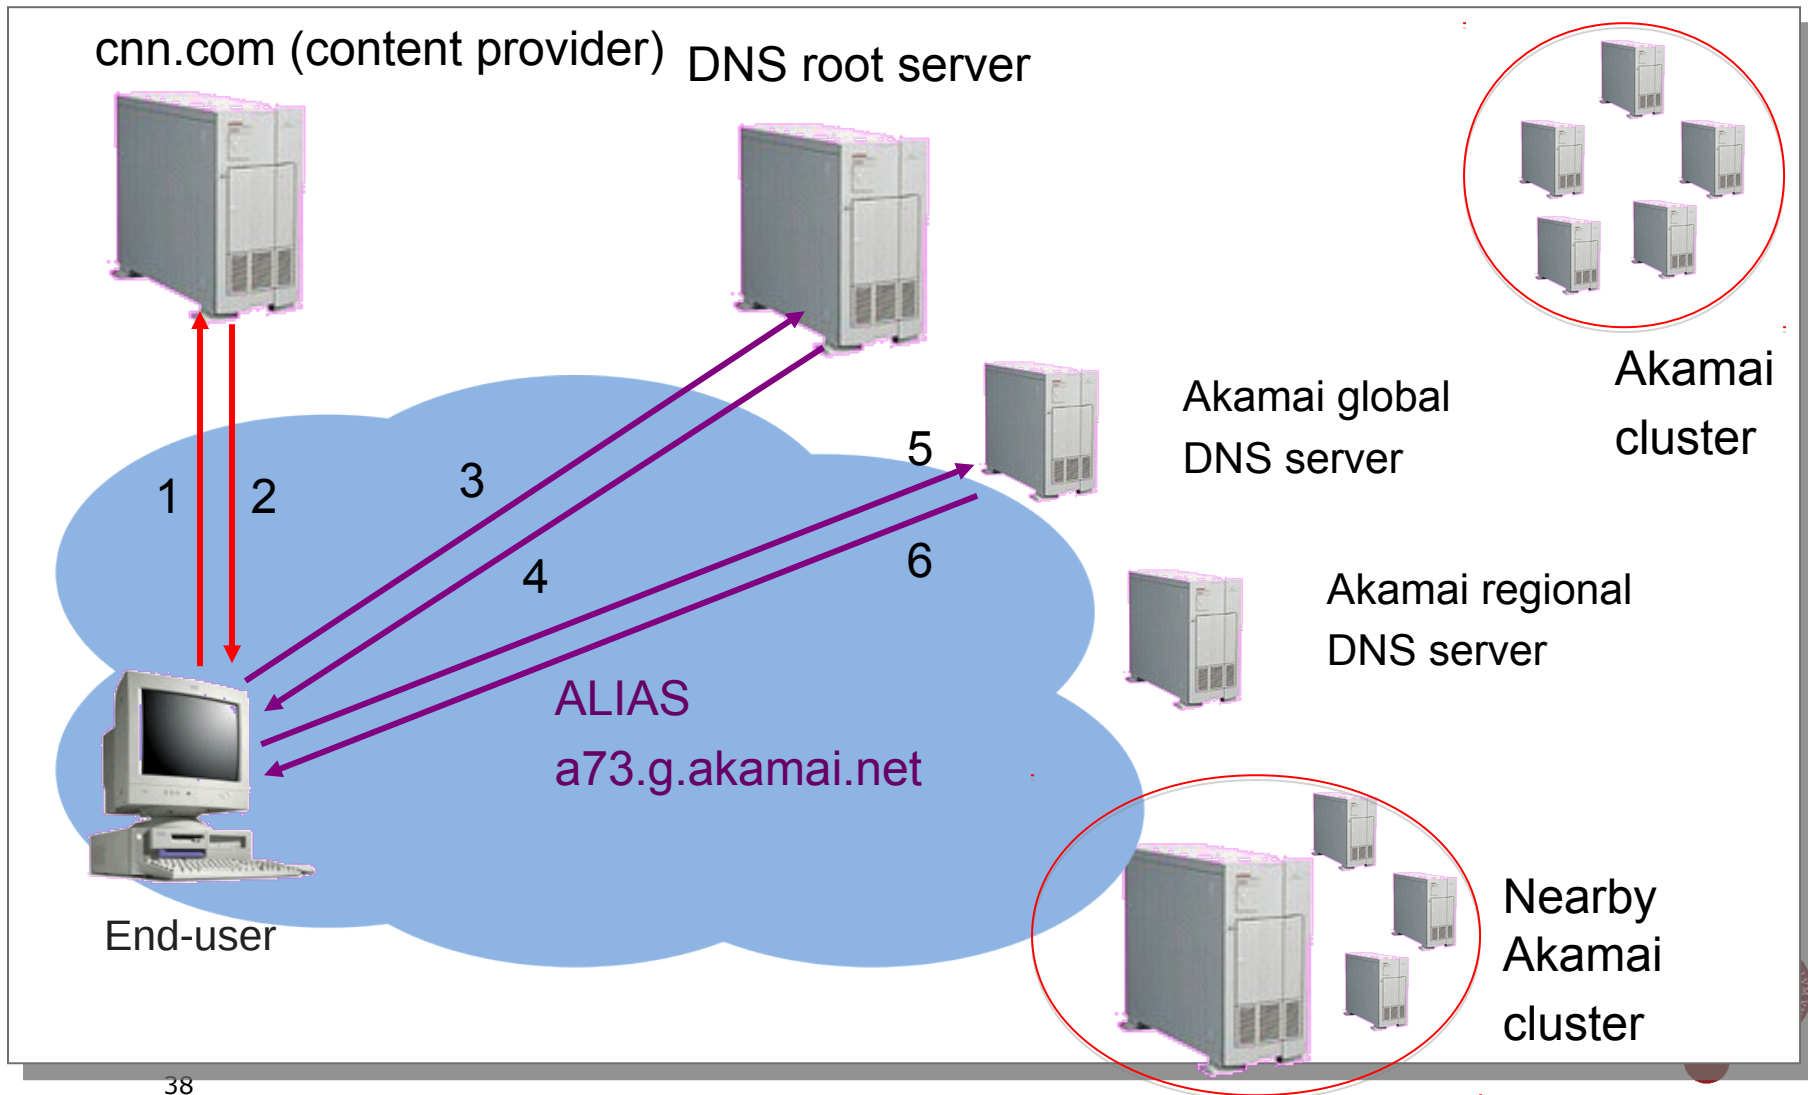
How Akamai Works

Source: Freedman



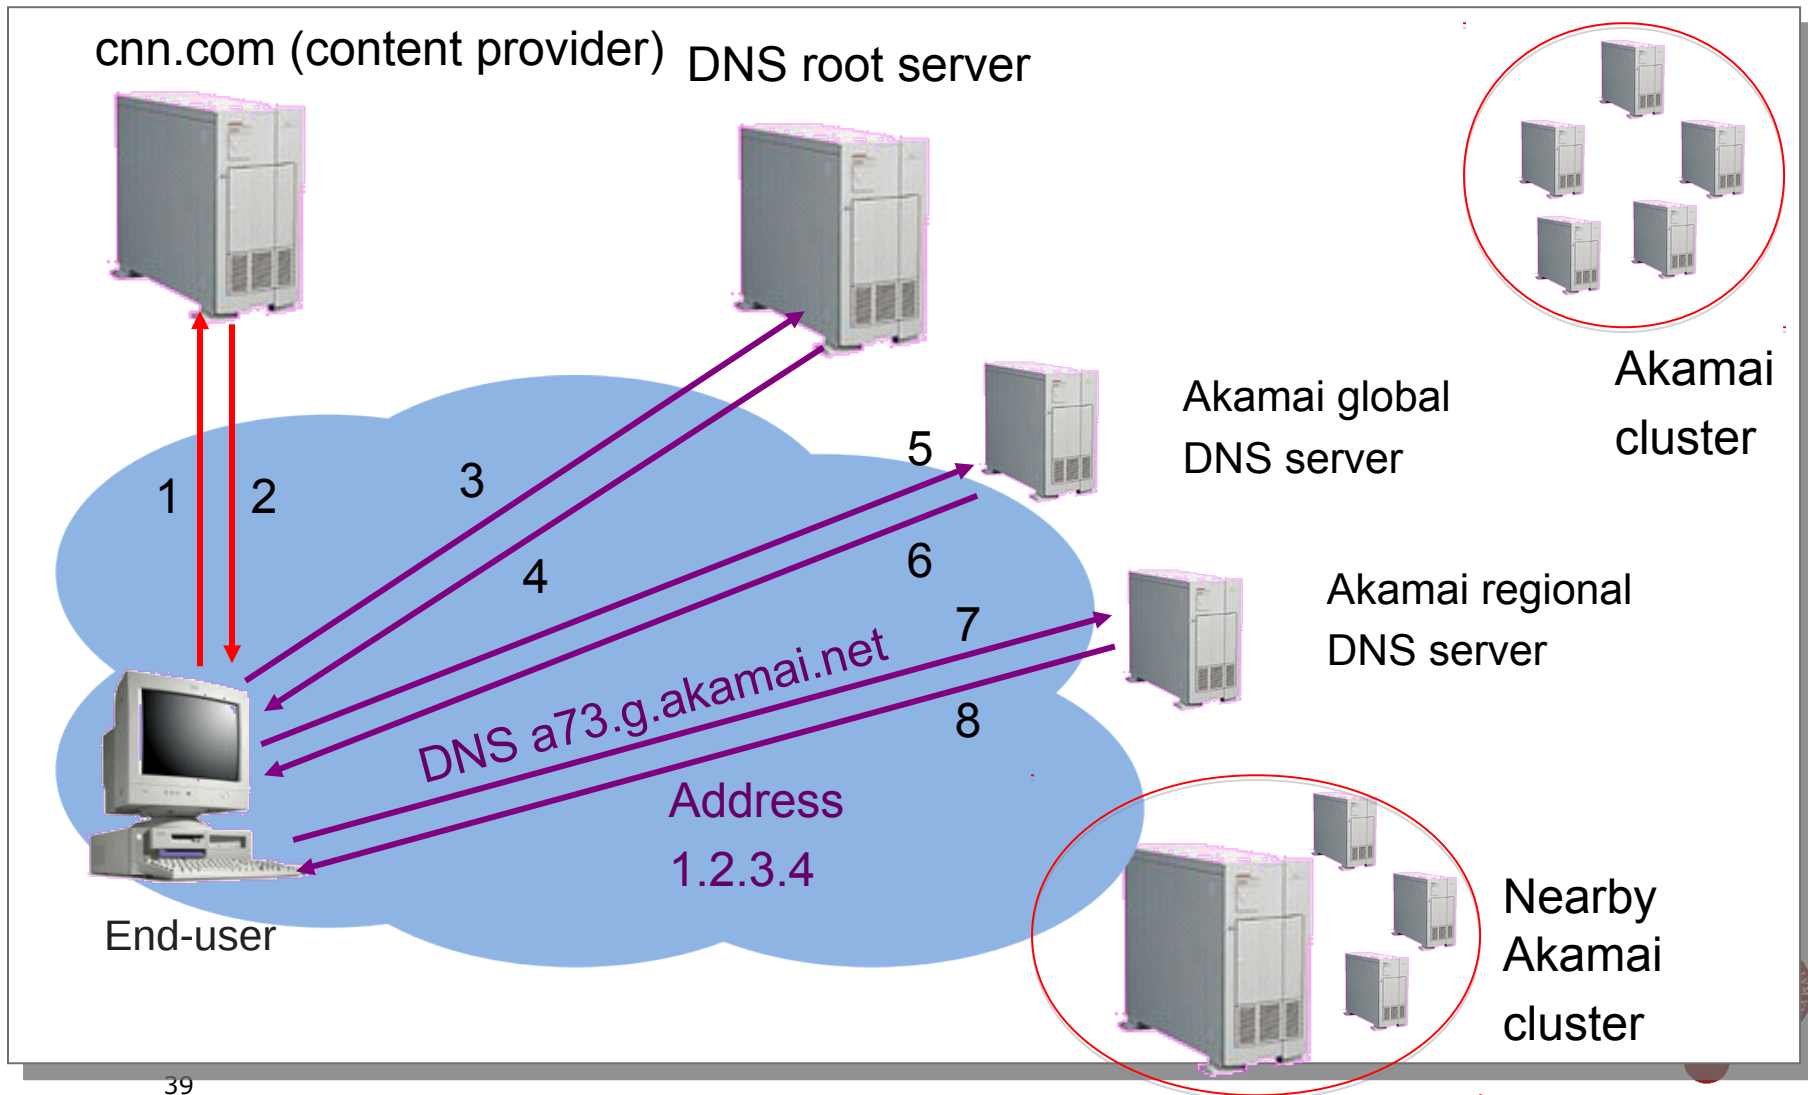
How Akamai Works

Source: Freedman



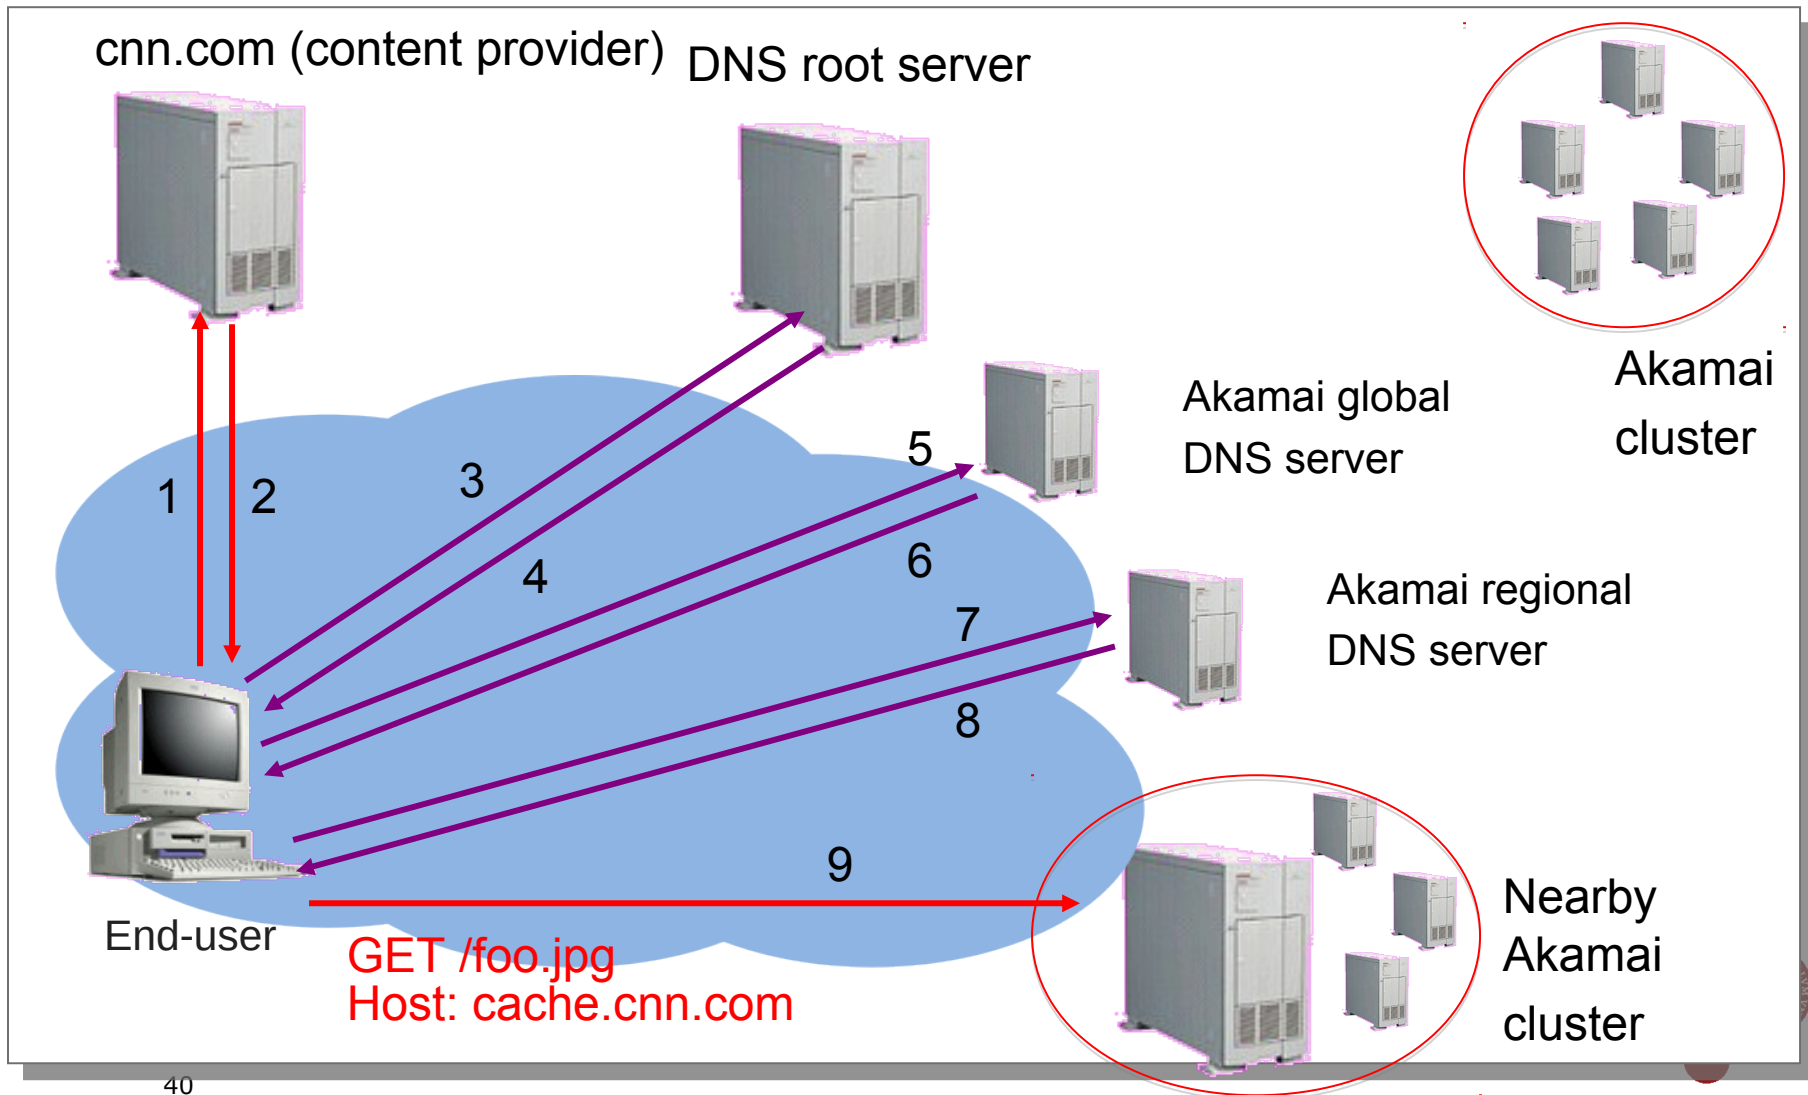
How Akamai Works

Source: Freedman



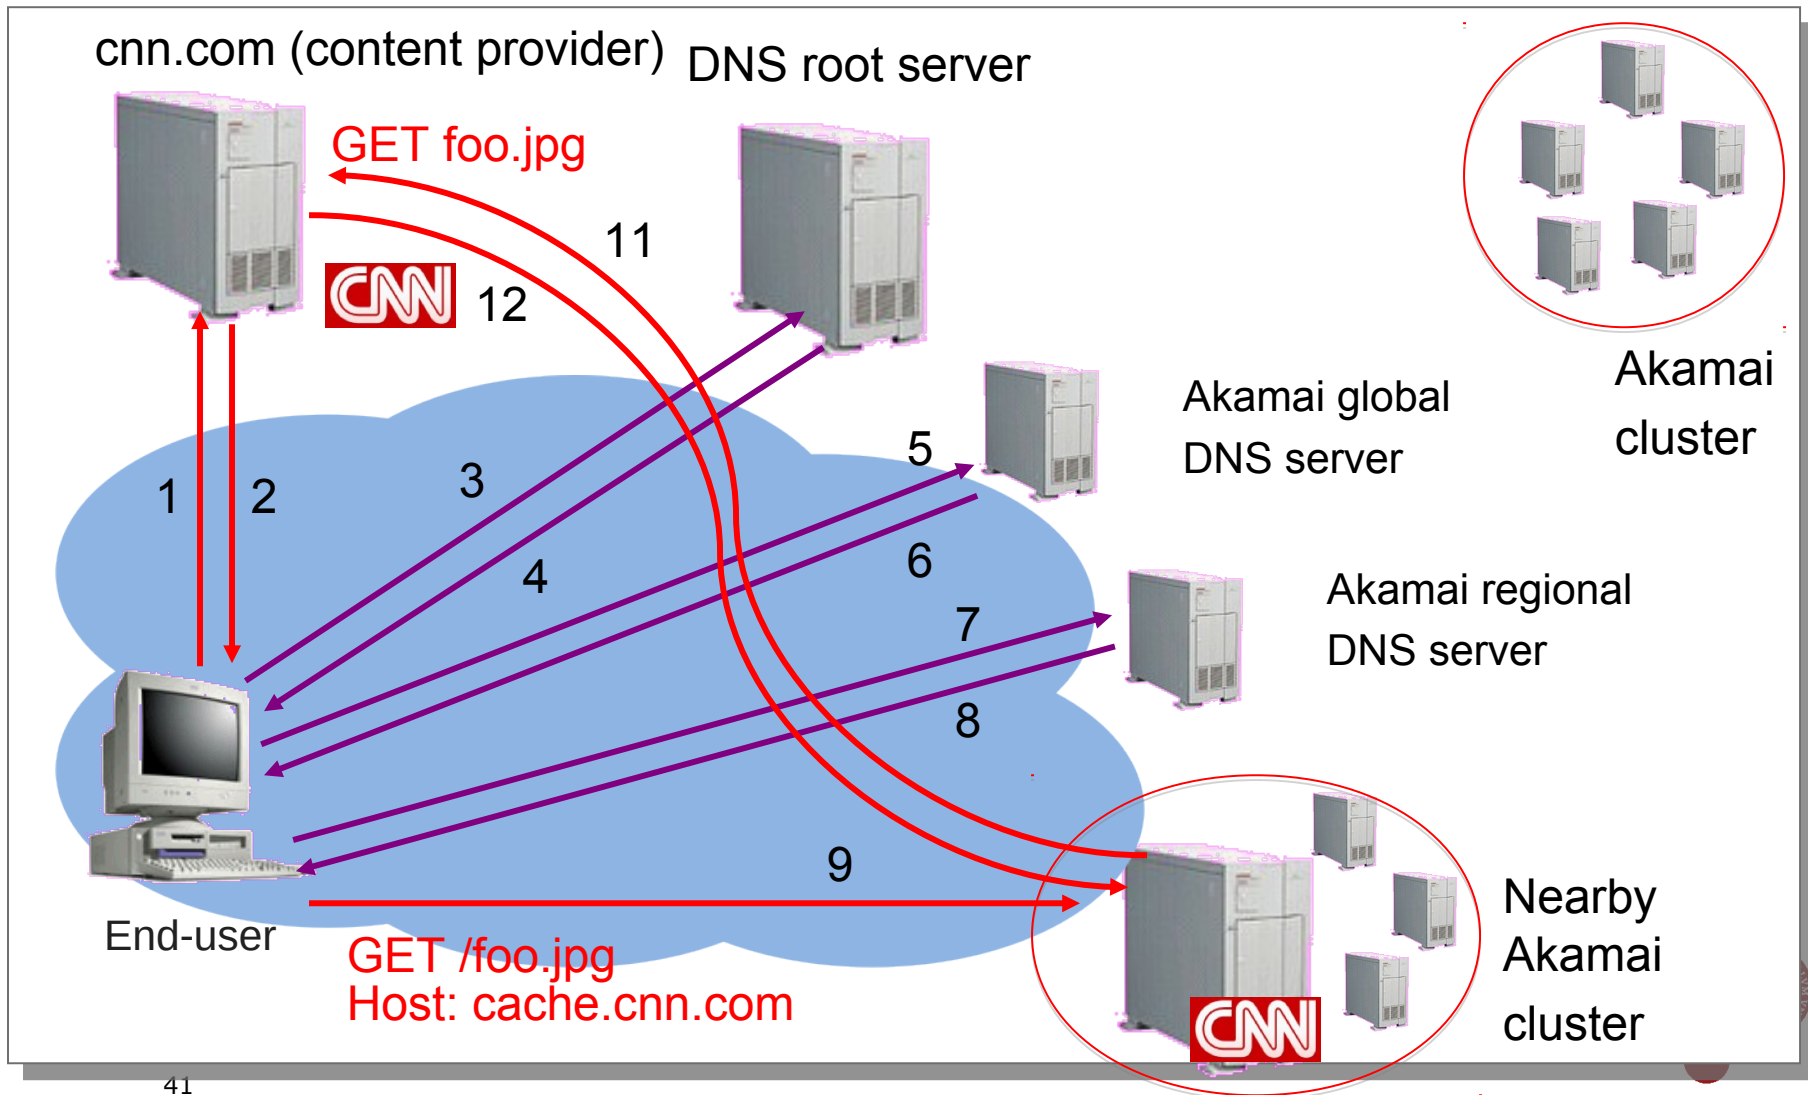
How Akamai Works

Source: Freedman



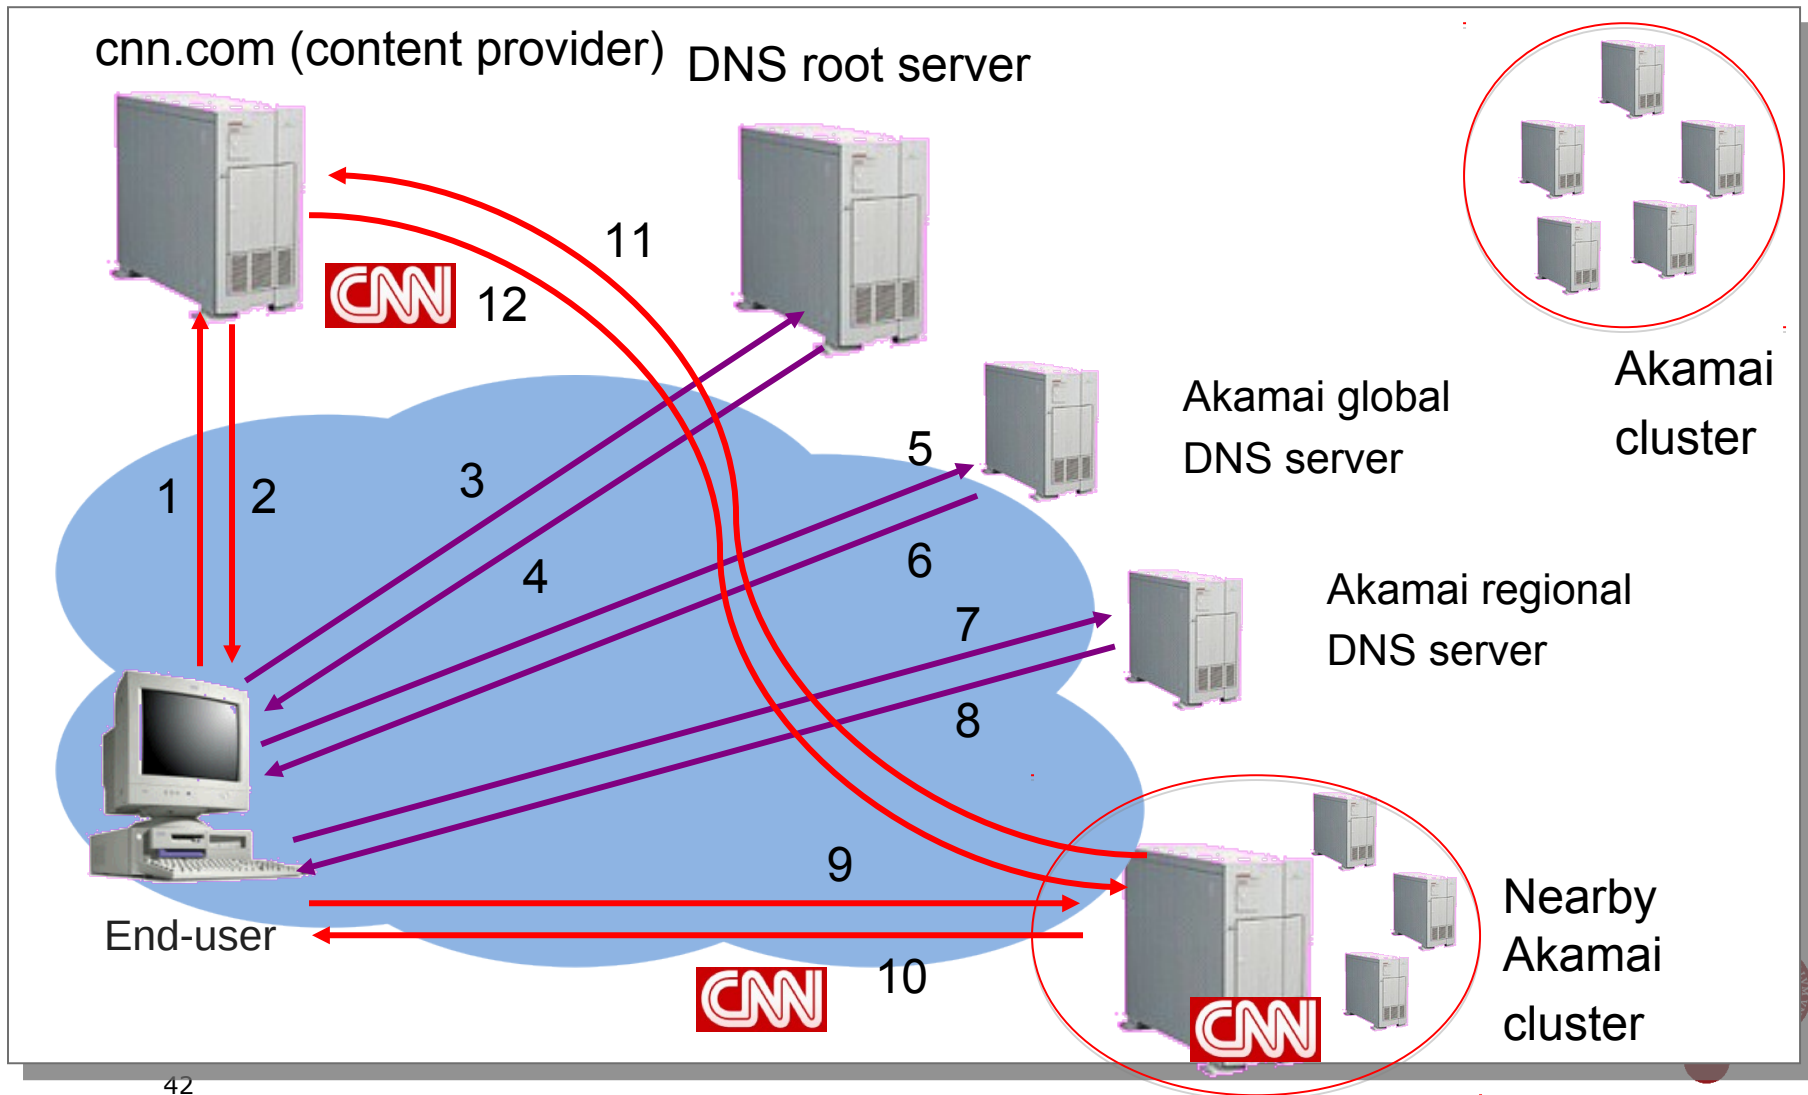
How Akamai Works

Source: Freedman



How Akamai Works

Source: Freedman

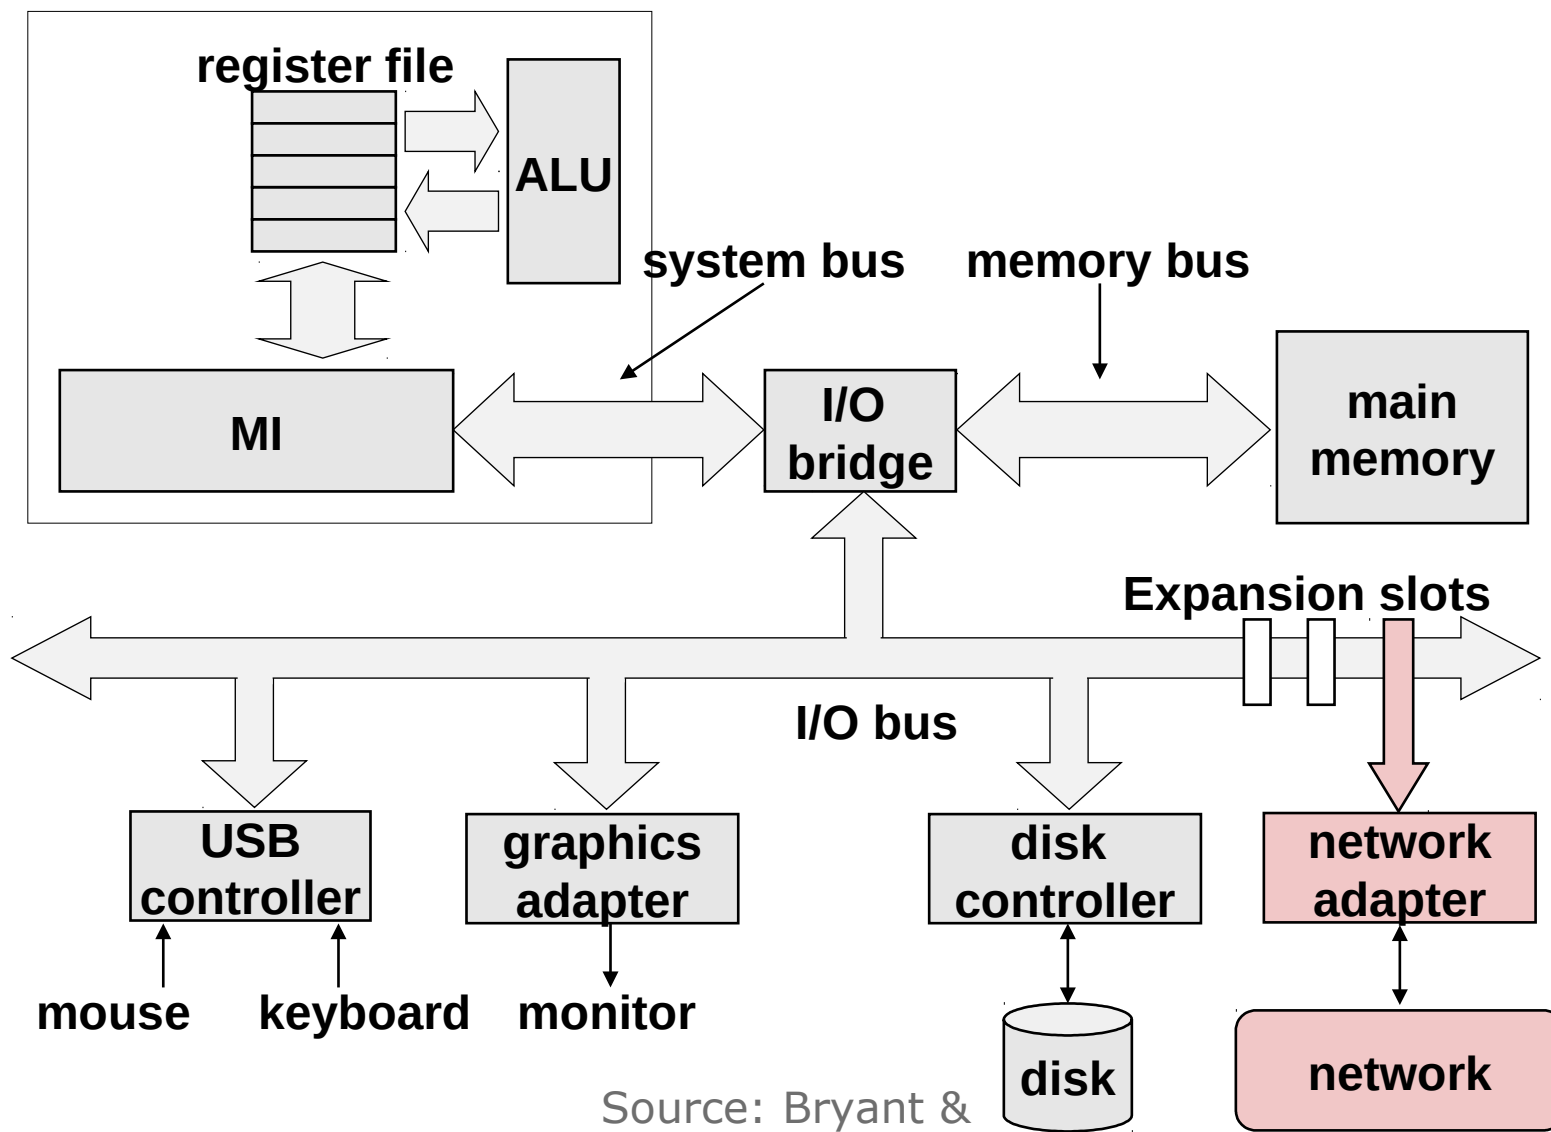


Summary

- Network applications
 - Email, Web → more in textbook, Chapter 2!
- Socket abstraction
 - Communication between processes
 - Client / Server, Peer-to-Peer
- HTTP concepts
 - Web objects, request / response (pull)
 - Persistent connections, web proxies
 - Caching, content delivery
- Assignment A6 has been released today



What's next? More network programming in C ?



Source: Bryant &
Halloran

