




Network layer: DHCP, NAT, Tunneling & Routing

Vivek Shah

Based on slides compiled by
Marcos Vaz Salles



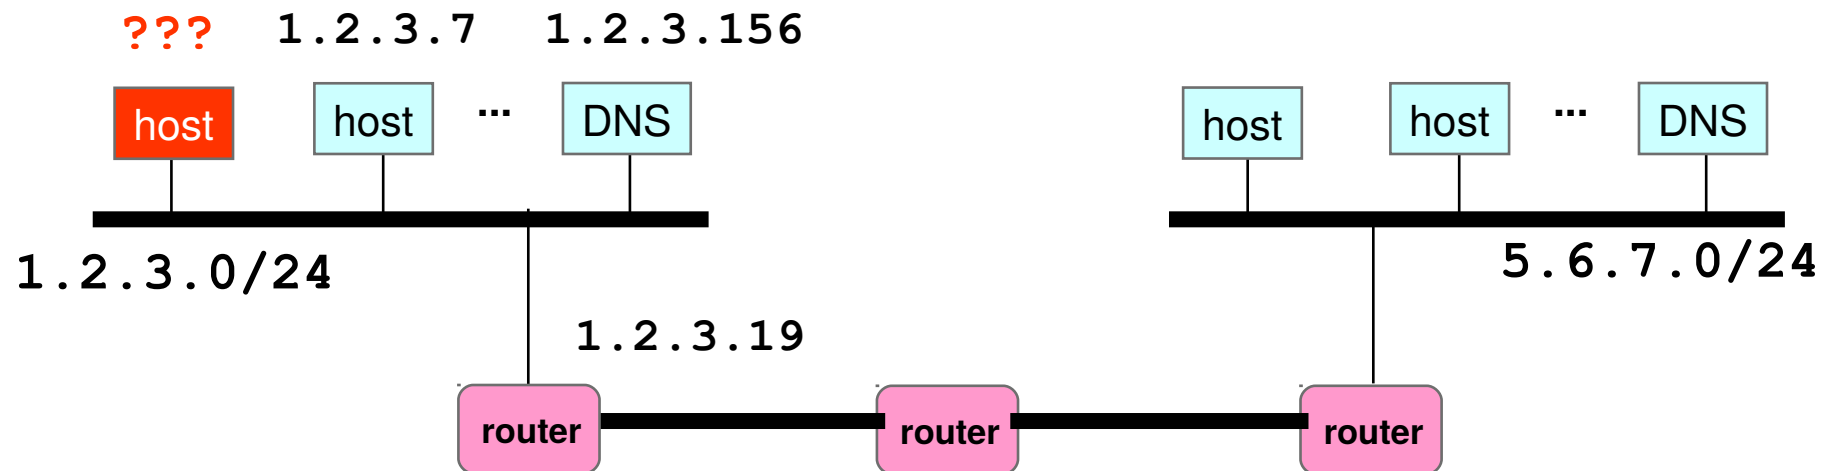
Recap: Network layer

- What is the “best effort guarantee” of network layer ?
- Why can fragmentation happen at routers ?
How does IPv4 handle it ? Why does IPv6 not handle it ?
- What do forwarding tables in routers contain ?
Why is longest prefix match chosen ?
- Why is an IP address hierarchical ? What is a subnet mask ?
- How are IP addresses allocated ?



How To Bootstrap an End Host?

- What local Domain Name System server to use?
- What IP address the host should use?
- How to send packets to remote destinations?
- How to ensure incoming packets arrive?

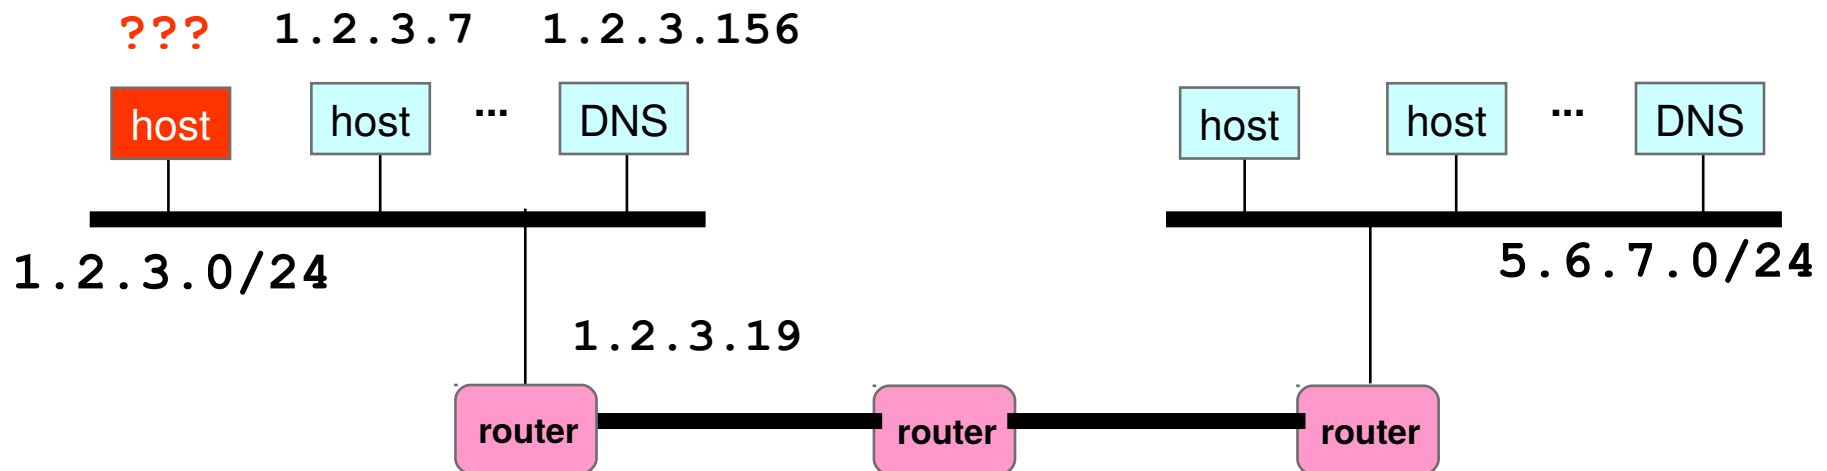


Source: Freedman



Avoiding Manual Configuration

- Dynamic Host Configuration Protocol (DHCP)
 - End host learns how to send packets
 - Learn IP address, DNS servers, and gateway
- Address Resolution Protocol (ARP)
 - Others learn how to send packets to the end host
 - Learn mapping between IP address & interface address



Source: Freedman




Key Ideas in Both Protocols

- **Broadcasting:** when in doubt, shout!
 - Broadcast query to all hosts in the local-area-network
 - ... when you don't know how to identify the right one
- **Caching:** remember the past for a while
 - Store the information you learn to reduce overhead
 - Remember your own address & other host's addresses
- **Soft state:** ... but eventually forget the past
 - Associate a time-to-live field with the information
 - ... and either refresh or discard the information
 - Key for robustness in the face of unpredictable change




Media Access Control (MAC) Addresses



Source: Freedman

Bootstrapping Problem

- Host doesn't have an IP address yet
 - So, host doesn't know what source address to use
- Host doesn't know who to ask for an IP address
 - So, host doesn't know what destination addr to use
- Solution: shout to discover a server who can help
 - Broadcast a DHCP server-discovery message
 - Server sends a DHCP "offer" offering an address

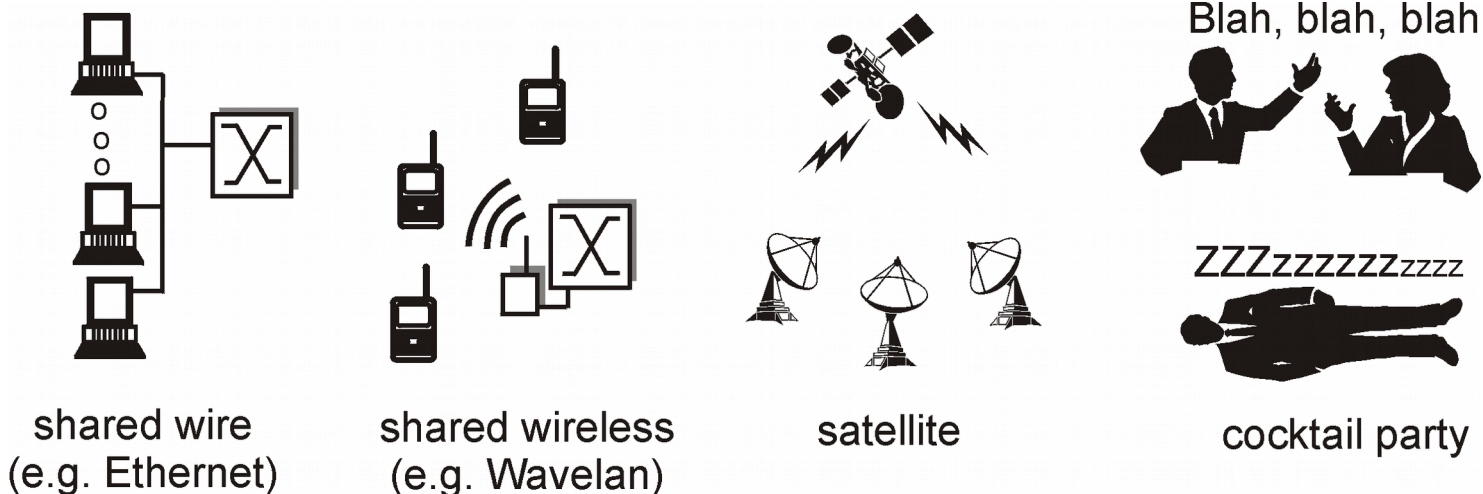


Source: Freedman



Broadcasting

- Broadcasting: sending to everyone
 - Special destination address: FF-FF-FF-FF-FF-FF
 - All adapters on the LAN receive the packet
- Delivering a broadcast packet
 - Easy on a “shared media”
 - Like shouting in a room – everyone can hear you



Source: Freedman

Response from the DHCP Server

- DHCP “offer message” from the server
 - Configuration parameters (proposed IP address, mask, gateway router, DNS server, ...)
 - Lease time (the time the information remains valid)
- Multiple servers may respond
 - Multiple servers on the same broadcast media
 - Each may respond with an offer
 - The client can decide which offer to accept
- Accepting one of the offers
 - Client sends a DHCP request echoing the parameters
 - The DHCP server responds with an ACK to confirm
 - ... and the other servers see they were not chosen



DHCP client-server scenario

DHCP server: 223.1.2.5

DHCP discover

src : 0.0.0.0, 68
dest.: 255.255.255.255, 67
yiaddr: 0.0.0.0
transaction ID: 654

arriving
client



DHCP offer

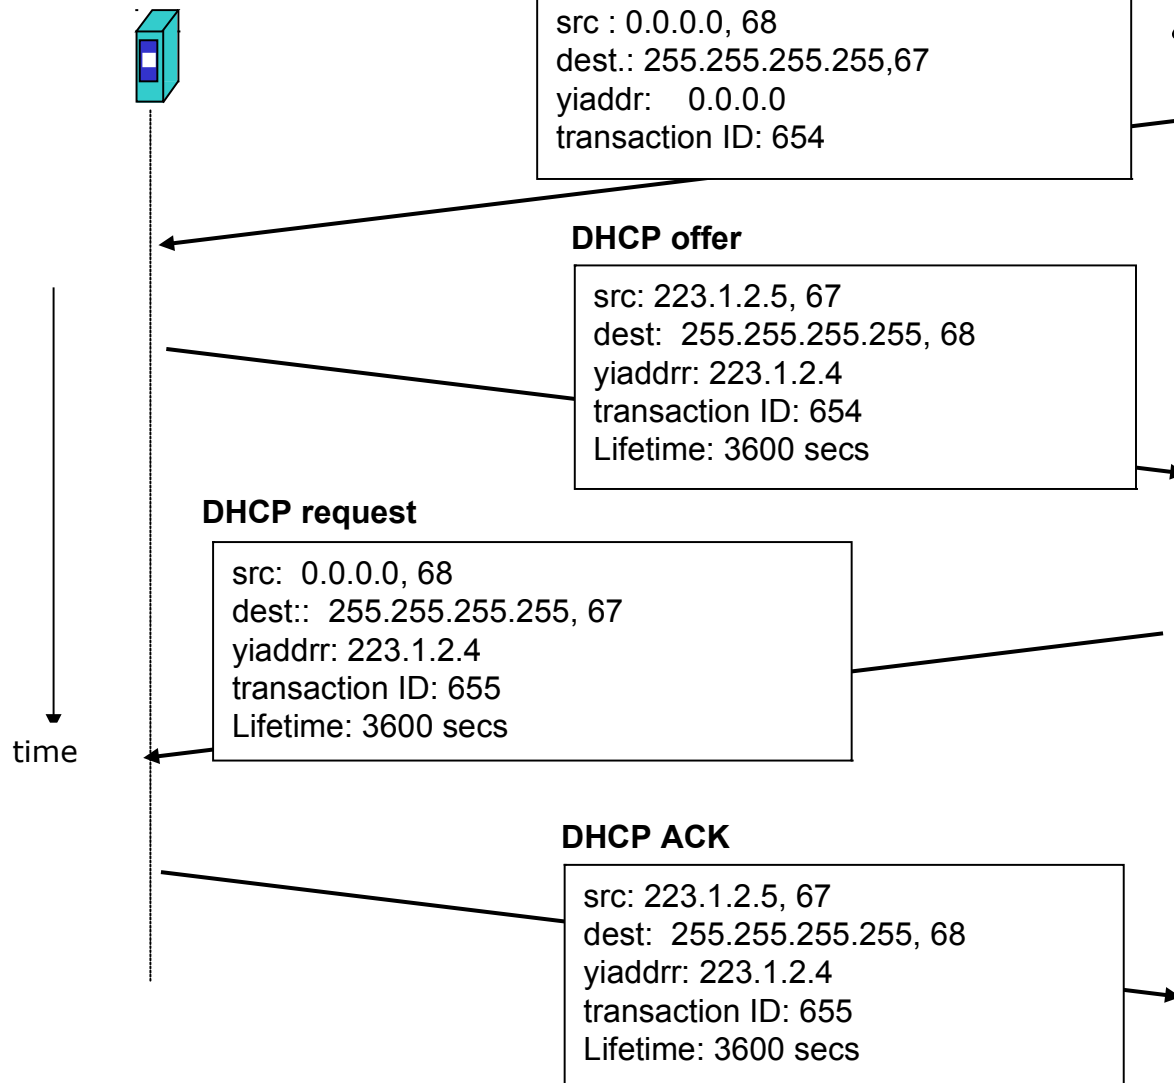
src: 223.1.2.5, 67
dest: 255.255.255.255, 68
yiaddr: 223.1.2.4
transaction ID: 654
Lifetime: 3600 secs

DHCP request

src: 0.0.0.0, 68
dest.: 255.255.255.255, 67
yiaddr: 223.1.2.4
transaction ID: 655
Lifetime: 3600 secs

DHCP ACK

src: 223.1.2.5, 67
dest: 255.255.255.255, 68
yiaddr: 223.1.2.4
transaction ID: 655
Lifetime: 3600 secs



Why are DHCP
request and ACK
broadcast ?

Source:
Kurose &
Ross

Deciding What IP Address to Offer

- Server as centralized configuration database
 - All parameters are statically configured in the server
 - E.g., a dedicated IP address for each MAC address
 - Avoids complexity of configuring hosts directly
 - ... while still having a permanent IP address per host
- Or, dynamic assignment of IP addresses
 - Server maintains a pool of available addresses
 - ... and assigns them to hosts on demand
 - Leads to less configuration complexity
 - ... and more efficient use of the pool of addresses
 - Though, it is harder to track the same host over time



Soft State: Refresh or Forget

- Why is a lease time necessary?
 - Client can release the IP address (DHCP RELEASE)
 - E.g., "ipconfig /release" at the DOS prompt
 - E.g., clean shutdown of the computer
 - But, the host might not release the address
 - E.g., the host crashes (blue screen of death!)
 - E.g., buggy client software
 - And you don't want the address to be allocated forever
- Performance trade-offs
 - Short lease time: returns inactive addresses quickly
 - Long lease time: avoids overhead of frequent renewals



Middleboxes

- Middleboxes are intermediaries
 - Interposed in-between the communicating parties
 - Often without knowledge of original data
- Myriad uses
 - Network address translators
 - Firewalls
 - Tunnel endpoints
 - Traffic shapers
 - Intrusion detection systems
 - Transparent Web proxy caches
 - Application accelerators

“An abomination!”

- Violation of layering
- Hard to reason about
- Responsible for subtle bugs

“A practical necessity!”

- Solve real/pressing problems
- Needs not likely to go away

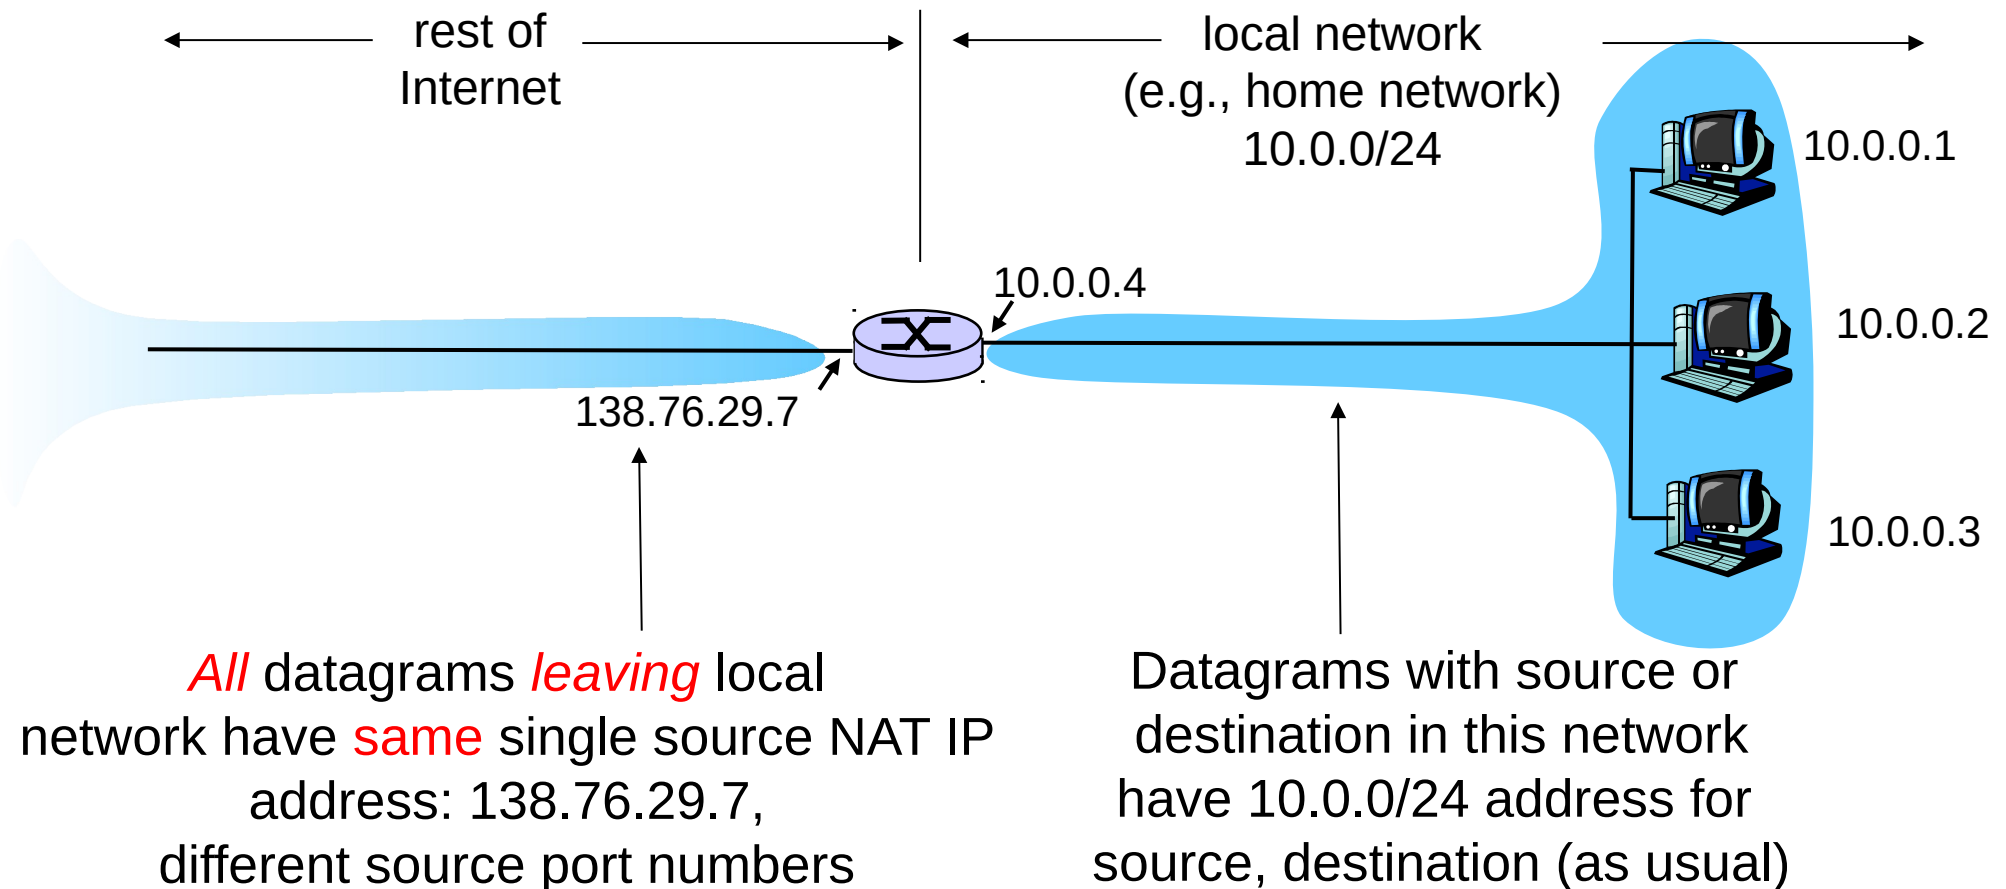


History of NATs

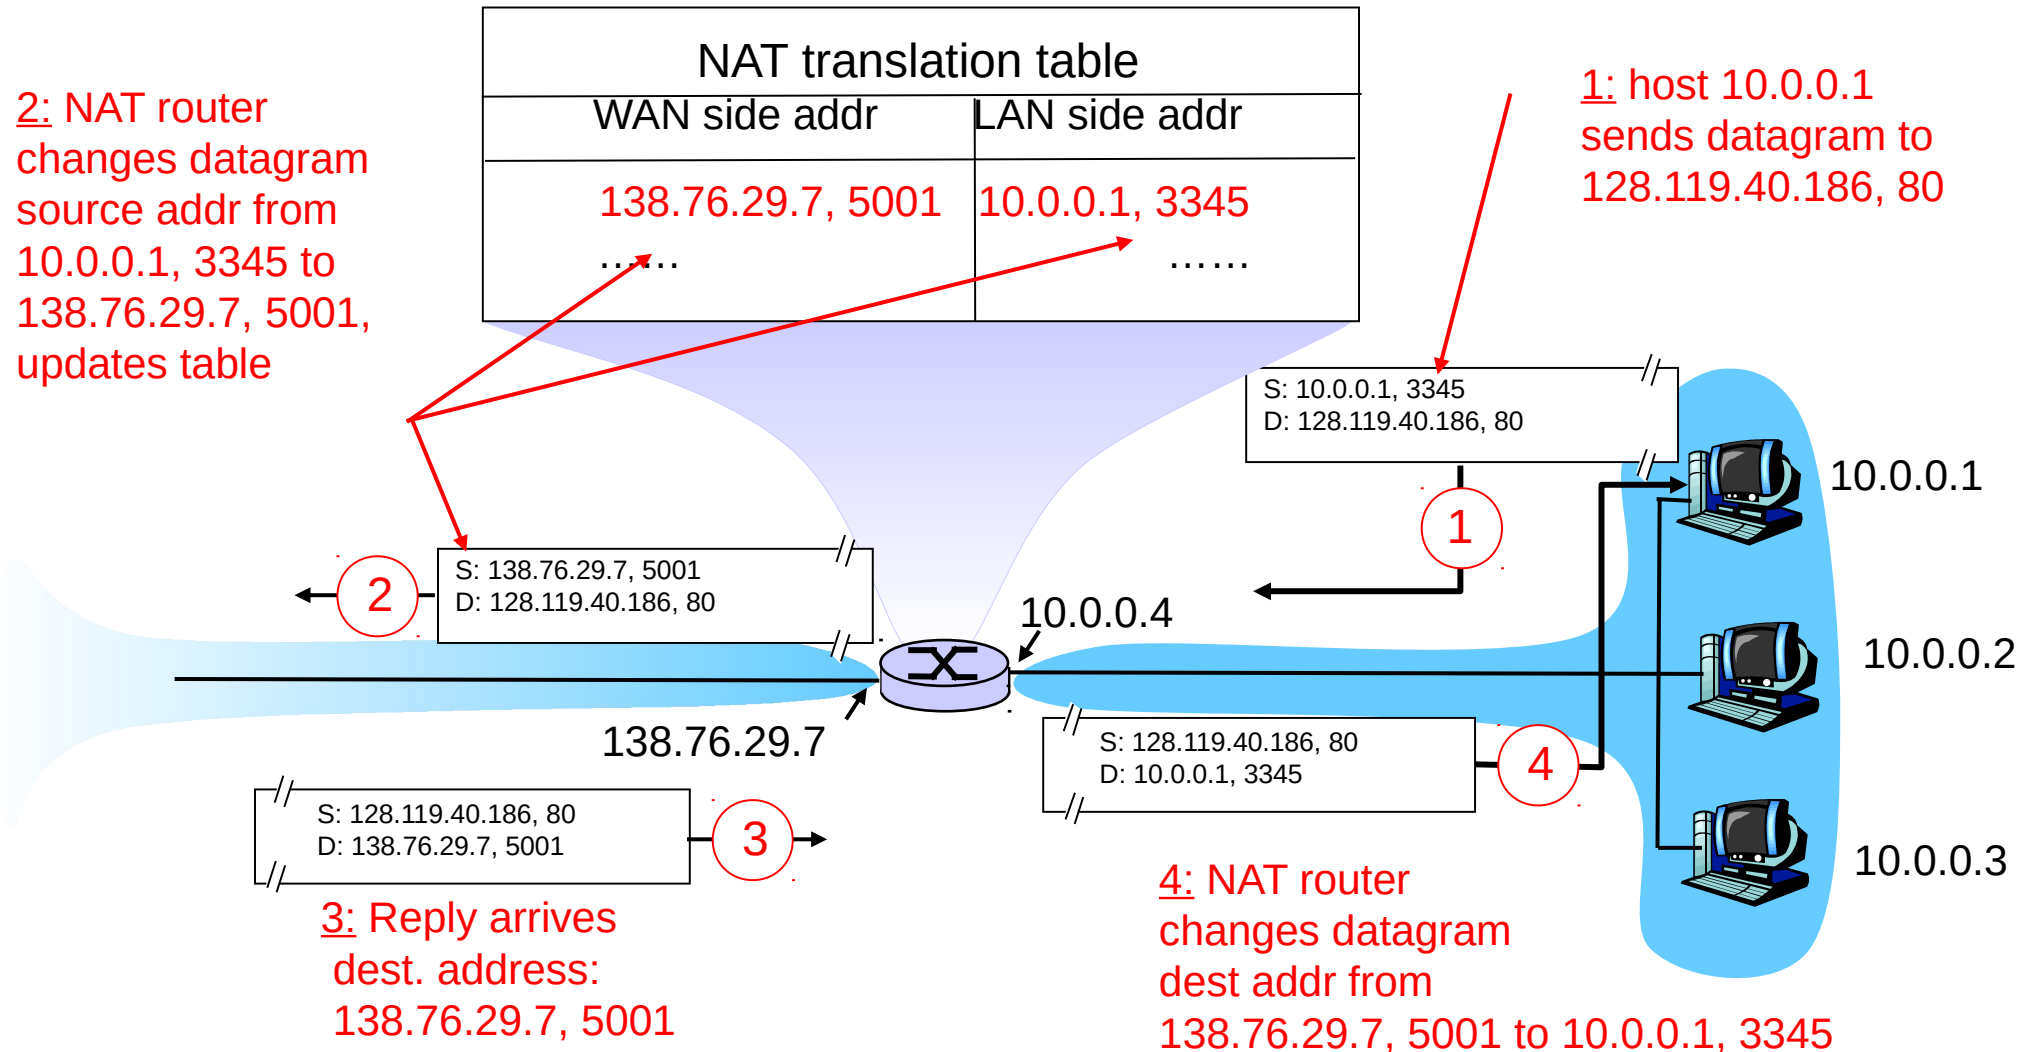
- IP address space depletion
 - Clear in early 90s that 2^{32} addresses not enough
 - Work began on a successor to IPv4
- In the meantime...
 - Share addresses among numerous devices
 - ... without requiring changes to existing hosts
- Meant to provide short-term remedy
 - Now: NAT is widely deployed, much more than IPv6



NAT: Network Address Translation



NAT: Network Address Translation



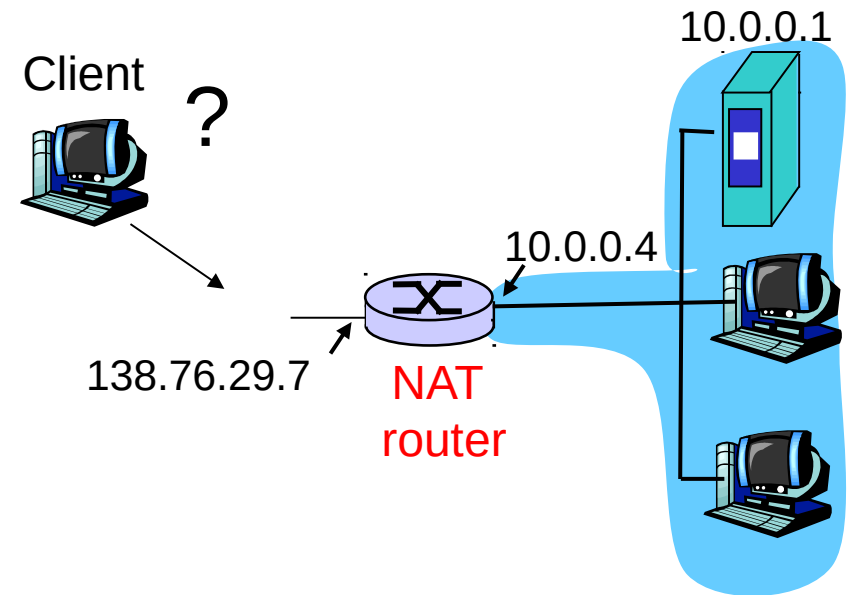
Maintaining the Mapping Table

- Create an entry upon seeing an outgoing packet
 - Packet with new (source addr, source port) pair
- Eventually, need to delete entries to free up #'s
 - When? If no packets arrive before a timeout
 - (At risk of disrupting a temporarily idle connection)
- Yet another example of “soft state”
 - I.e., removing state if not refreshed for a while



NAT Traversal Problem

- client wants to connect to server with address 10.0.0.1
 - server address 10.0.0.1 local to LAN (client can't use it as destination addr)
 - only one externally visible NATed address: 138.76.29.7



- How can we deal with incoming connections?
- How do we map to multiple services inside the NAT'ed subnet?

Where is NAT Implemented?

- Home router (e.g., Linksys box)
 - Integrates router, DHCP server, NAT, etc.
 - Use single IP address from the service provider
- Campus or corporate network
 - NAT at the connection to the Internet
 - Share a collection of public IP addresses
 - Avoid complexity of renumbering hosts/routers when changing ISP (w/ provider-allocated IP prefix)



NAT limitations and criticism

- 16-bit port-number field:
 - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
 - routers should only process up to layer 3
 - violates end-to-end argument
 - NAT possibility must be taken into account by app designers, e.g., P2P applications
 - address shortage should instead be solved by IPv6




Transition From IPv4 To IPv6

- Not all routers can be upgraded simultaneous
 - no “flag days”
 - How will the network operate with mixed IPv4 and IPv6 routers?
- *Tunneling*: IPv6 carried as payload in IPv4 datagram among IPv4 routers



IP Tunneling to Build Overlay Links

- IP tunnel is a virtual point-to-point link
 - Illusion of a direct link between two separated nodes



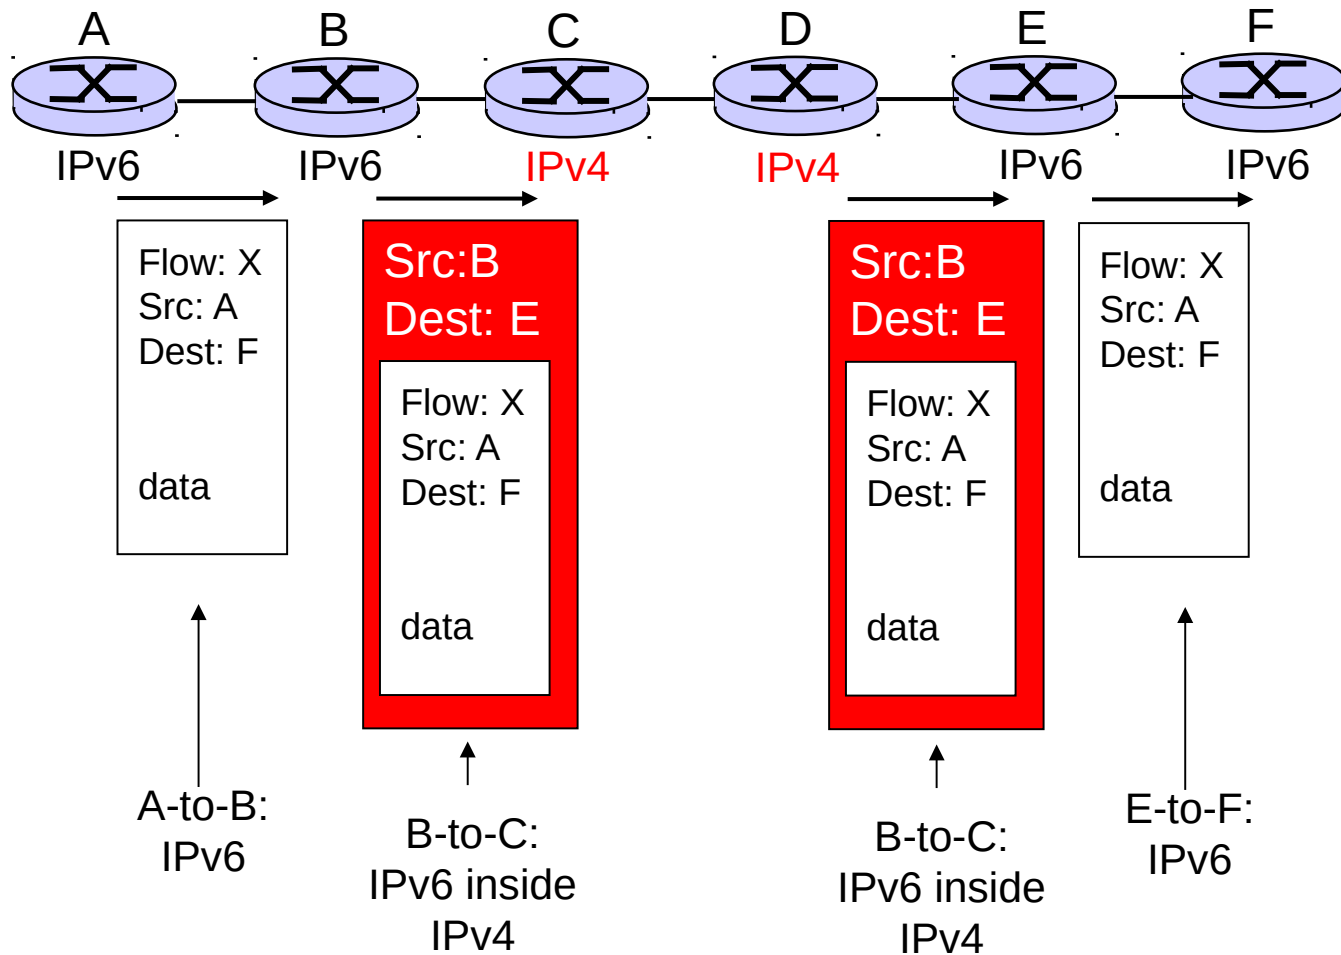
- Encapsulation of the packet inside an IP datagram
 - Node B sends a packet to node E
 - ... containing another packet as the payload

IP Tunneling

Logical view:




Physical view:



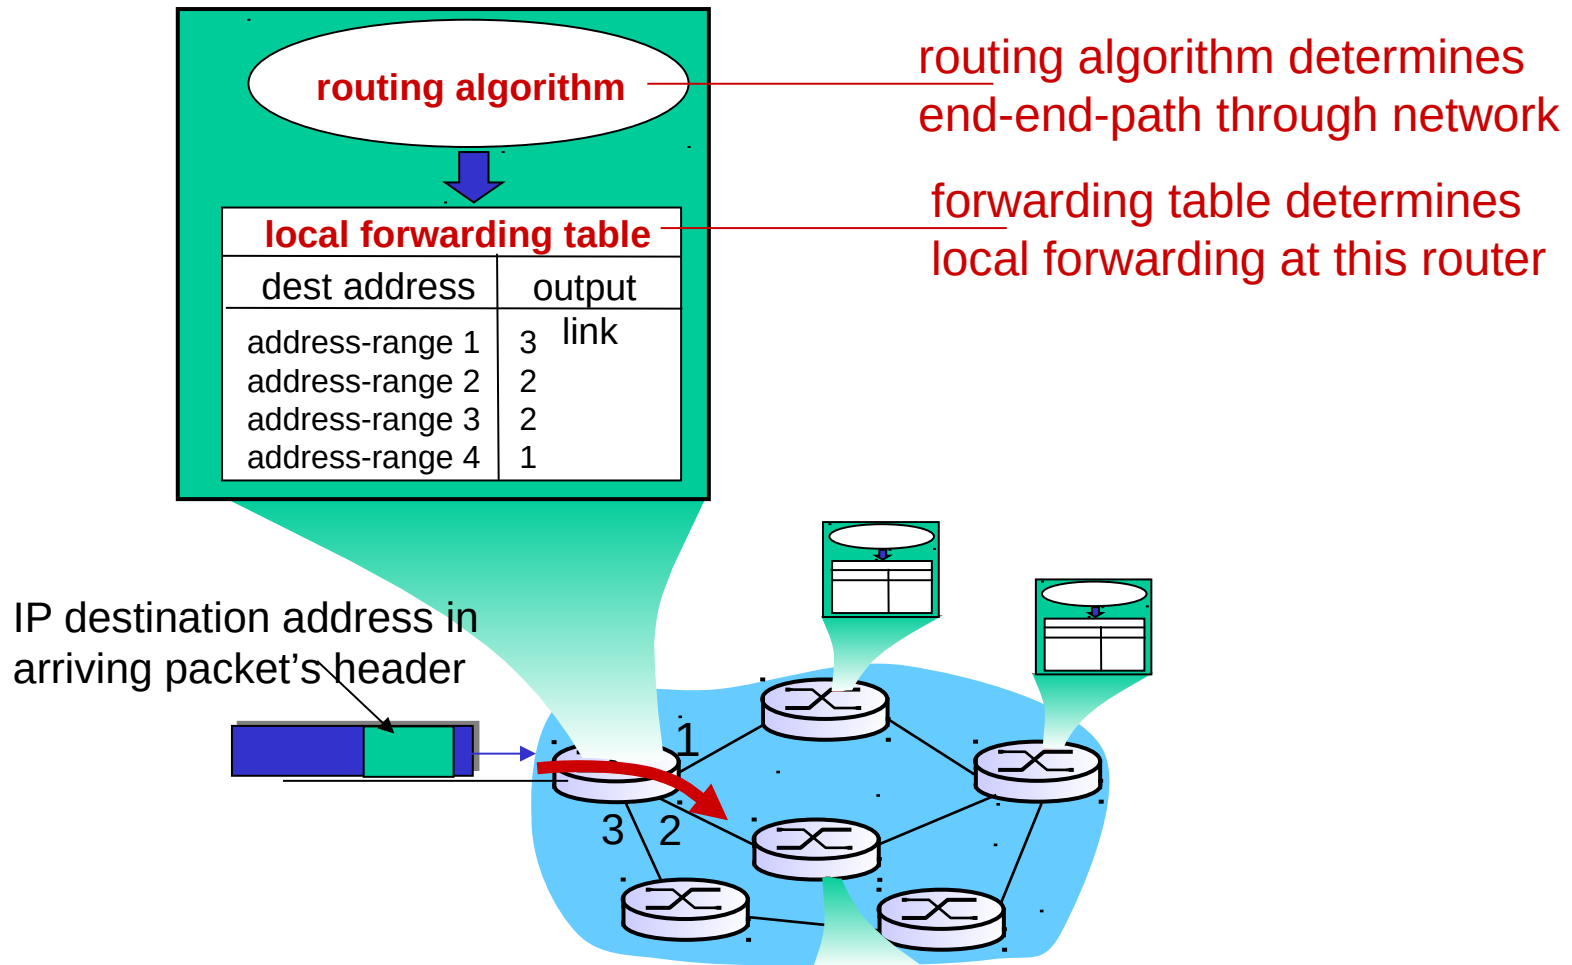
Source:
Kurose
& Ross



Routing: Mapping Link to Path



Interplay of routing and forwarding



Source:
Kurose
& Ross

Three Issues to Address

- What does the protocol compute?
 - E.g., shortest paths
- What algorithm does the protocol run?
 - E.g., link-state routing
- How do routers learn end-host locations?
 - E.g., injecting into the routing protocol



Routing Algorithm Classification

Q: global or decentralized information?

global:

- all routers have complete topology, link cost info
- “link state” algorithms

decentralized:

- router knows physically-connected neighbors, link costs to neighbors
- iterative process of computation, exchange of info with neighbors
- “distance vector” algorithms

Q: static or dynamic?

static:

- routes change slowly over time

dynamic:


- routes change more quickly
 - periodic update
 - in response to link cost changes

Source: Kurose & Ross




What to Compute ?

- Shortest path(s) between pairs of nodes
 - A shortest-path tree rooted at each node
 - Min hop count or min sum of edge weights



Shortest Path Problem

- Compute: *path costs to all nodes*
 - From a given source u to all other nodes
 - Cost of the path through each outgoing link
 - Next hop along the least-cost path to s



Link State : Dijkstra's Algorithm

- Flood the topology information to all nodes
- Each node computes shortest paths to other nodes

Initialization

```
S = {u}
for all nodes v
  if (v is adjacent to u)
    D(v) = c(u,v)
  else D(v) = ∞
```

Loop





```
add w with smallest D(w) to S
update D(v) for all adjacent v:
  D(v) = min{D(v), D(w) + c(w,v)}
until all nodes are in S
```

Used in OSPF and IS-IS

Source: Freedman







Link State : Routing Example



Source: Freedman



Link State : Routing Example (contd.)




Source: Freedman



Link State : Routing Example (contd.)

- Shortest-path tree from u
- Forwarding table at u



dest	link
v	(u,v)
w	(u,w)
x	(u,w)
y	(u,v)
z	(u,v)
s	(u,w)
t	(u,w)

Source: Freedman




Link State Algorithm Discussion

algorithm complexity: n nodes

- each iteration: need to check all nodes, w , not in N
- $n(n+1)/2$ comparisons: $O(n^2)$
- more efficient implementations possible: $O(n \log n)$

oscillations possible:


- e.g., support link cost equals amount of carried traffic:



Source: Kurose & Ross

Distance Vector : Bellman Ford Algorithm

- Define distances at each node x
 - $d_x(y)$ = cost of least-cost path from x to y
- Update distances based on neighbors
 - $d_x(y) = \min \{c(x,v) + d_v(y)\}$ over all neighbors v




$$d_u(z) = \min\{ c(u,v) + d_v(z), \\ c(u,w) + d_w(z) \}$$

Used in RIP and EIGRP


Source: Freedman

Distance Vector : Routing Example




$$d_y(z) = 1$$

$$d_x(z) = 4$$




$$d_v(z) = \min\{ 2+d_y(z), \\ 1+d_x(z) \} \\ = 3$$

Distance Vector : Routing Example (contd.)




$$d_w(z) = \min\{1+d_x(z), \\ 4+d_s(z), \\ 2+d_u(z) \} \\ = 5$$




$$d_u(z) = ??$$

(A) 5 (B) 6 (C) 7

Distance Vector : Routing Example (contd.)



$$d_w(z) = \min\{1+d_x(z), \\ 4+d_s(z), \\ 2+d_u(z) \} \\ = 5$$




$$d_u(z) = \min\{ 3+d_v(z), \\ 2+d_w(z) \} \\ = 6$$

Distance Vector : Link Cost Changes


link cost changes:

- Node detects local link cost change
- Updates routing info, recalculates distance vector
- If DV changes, notify neighbors
- **Good News travels fast**



link cost changes:

- Node detects local link cost change
 - **Bad news travels slow** - “count to infinity” problem!
 - 44 iterations before algorithm stabilizes: see text
 - **Poisoned Reverse for faster convergence.**
- Will this completely solve count to infinity problem?



Source: Kurose & Ross

Distance Vector : Link Cost Changes

message complexity

- **LS:** with n nodes, E links, $O(nE)$ msgs sent
- **DV:** exchange between neighbors only
 - convergence time varies

speed of convergence

- **LS:** $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- **DV:** convergence time varies
 - may be routing loops
 - count-to-infinity problem

robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its *own* table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate through network

Source: Kurose & Ross

Routing Issues

Our routing study thus far - idealization

- All routers identical
- Network “flat”
... *not* true in practice

Scale: with 600 million destinations:

- Can't store all dest's in routing tables!
- Routing table exchange would swamp links!

Administrative autonomy

- Internet = network of networks
- Each network admin may want to control routing in its own network

Source: Kurose & Ross

Hierarchical Routing – Standard CS trick

- Aggregate routers into regions, “autonomous systems” (AS)
- Routers in same AS run same routing protocol
 - “Intra-AS” routing protocol
 - Routers in different AS can run different intra-AS routing protocol

Gateway router:

- At “edge” of its own AS
- Has link to router in another AS

Source: Kurose & Ross



Summary

- DHCP: bootstrapping IP addresses
 - Broadcasting, caching, soft state
- NAT
 - A hack! ☺ Reading and reflection: Why?
- Many other hacks too! ☺
 - Tunneling, firewalls, mobile gateways, VPNs
- Routing Algorithms are graph based
 - Centralized → Link State
 - Distributed → Distance Vector
- Routing algorithms have different characteristics

